



GLitch: when WebGL turns into a hammer

Pietro Frigo

Cristiano Giuffrida, Herbert Bos, Kaveh Razavi

GLitch: what?

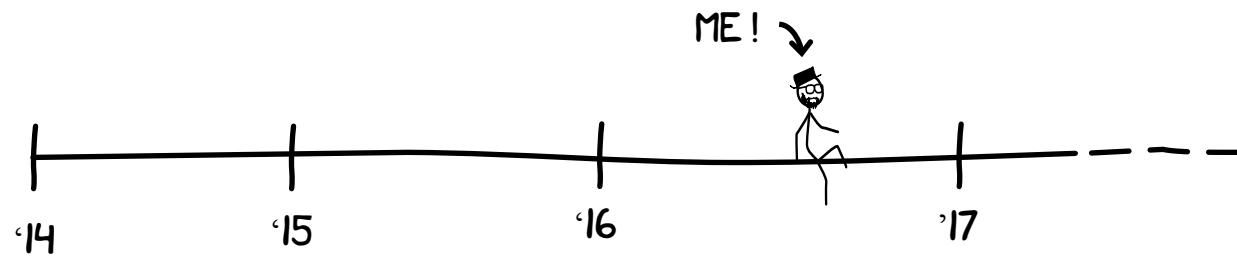
First Rowhammer exploit from JavaScript on mobile

Fastest JS-based Rowhammer exploit

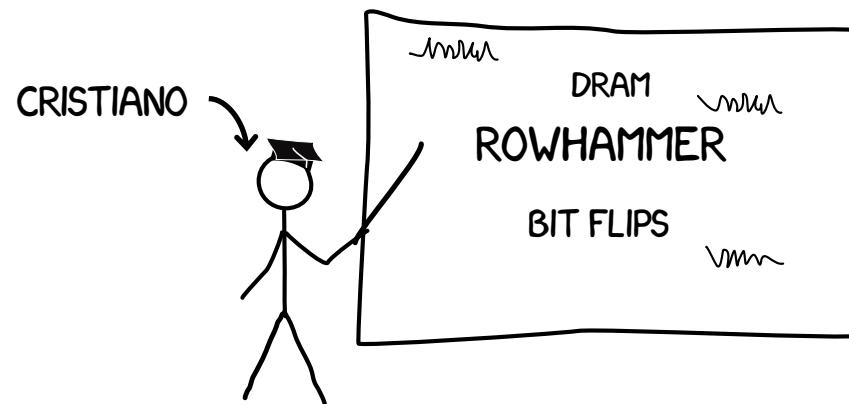
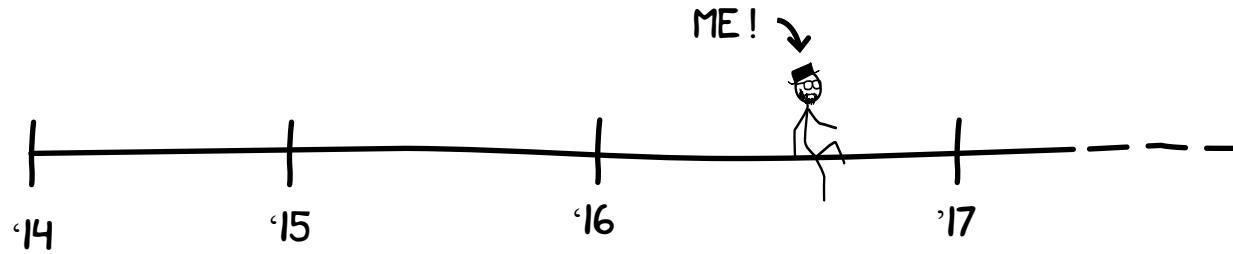
First GPU-accelerated bit flip

The chronicles of GLitch

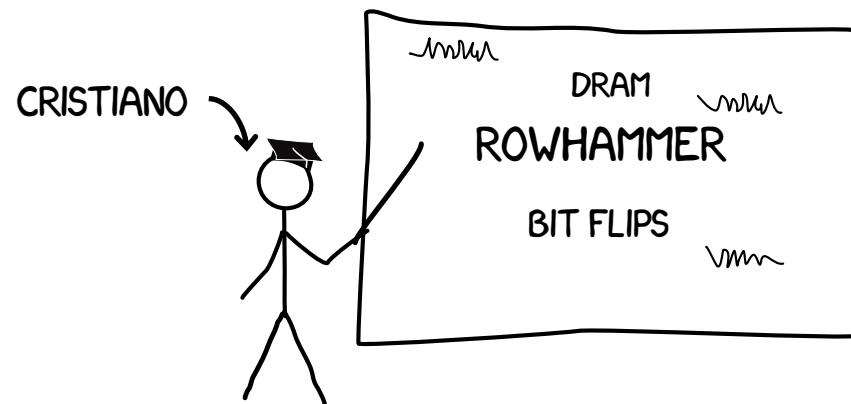
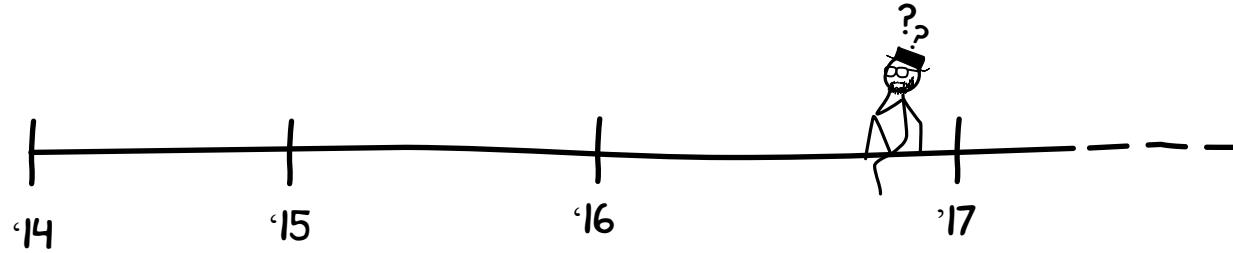
GLitch: the chronicles



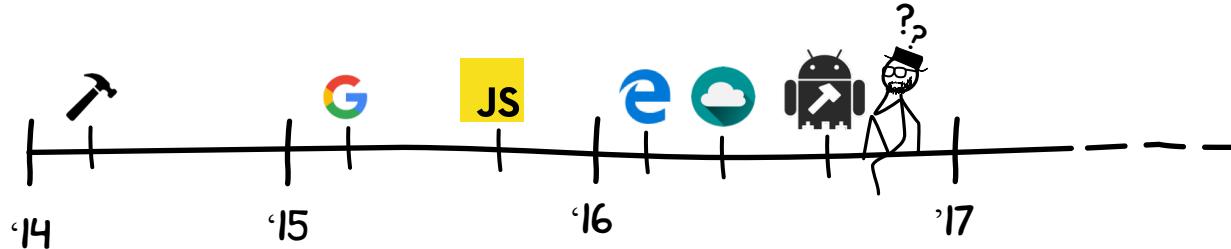
GLitch: the chronicles



GLitch: the chronicles



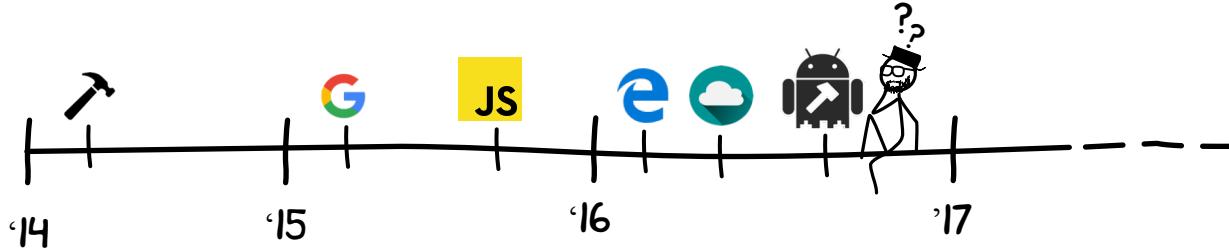
GLitch: the chronicles



- [1] Flipping bits in memory without accessing them
- [2] Google Project Zero: 1st Rowhammer root Exploit (flipping PTEs)
- [3] Rowhammer.js: 1st RH bit flip in JavaScript
- [4] Dedup est Machina: Breaking Microsoft Edge's sandbox
- [5] Flip Feng Shui: Breaking the cloud
- [6] Drammer: Flip feng shui goes mobile

everyween
+100XP

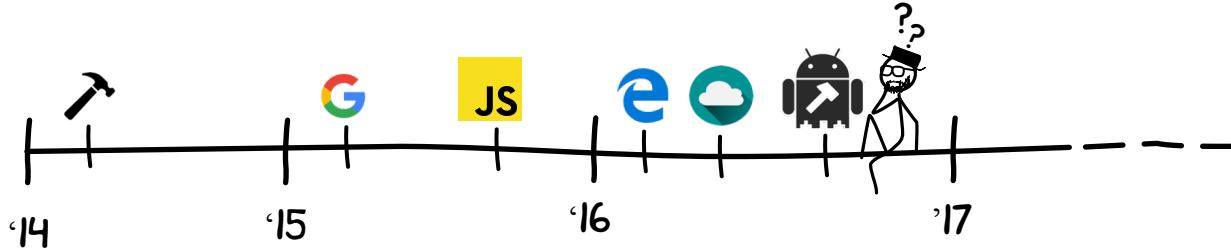
GLitch: the chronicles



- [1] Flipping bits in memory without accessing them
- [2] Google Project Zero: 1st Rowhammer root Exploit (flipping PTEs)
- [3] Rowhammer.js: 1st RH bit flip in JavaScript
- [4] Dedup est Machina: Breaking Microsoft Edge's sandbox
- [5] Flip Feng Shui: Breaking the cloud
- [6] Drammer: Flip feng shui goes mobile

everyween
+100XP

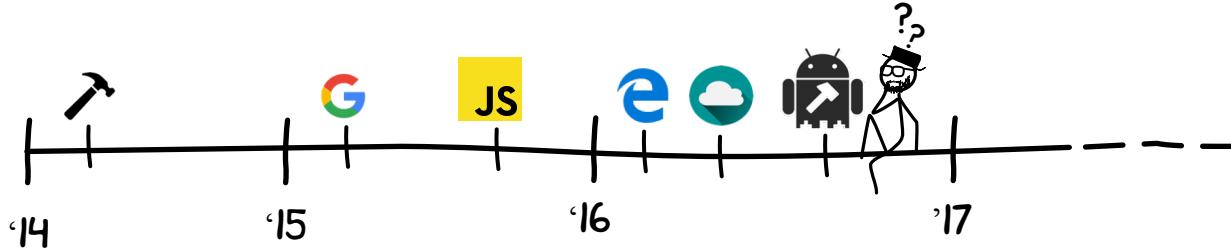
GLitch: the chronicles



- [1] Flipping bits in memory without accessing them
- [2] Google Project Zero: 1st Rowhammer root Exploit (flipping PTEs)
- [3] Rowhammer.js: 1st RH bit flip in JavaScript
- [4] Dedup est Machina: Breaking Microsoft Edge's sandbox
- [5] Flip Feng Shui: Breaking the cloud
- [6] Drammer: Flip feng shui goes mobile

everyween
+100XP

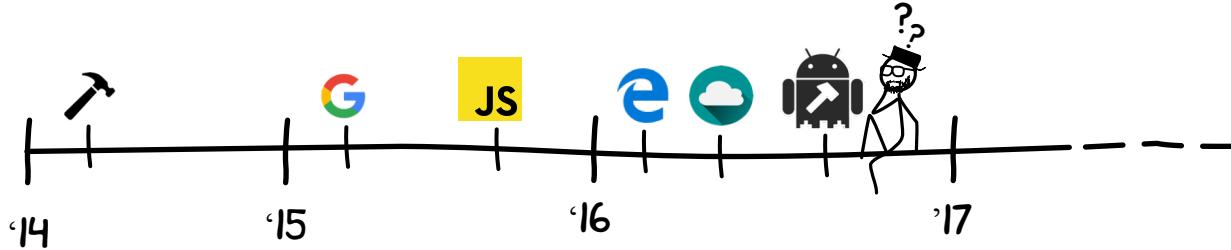
GLitch: the chronicles



- [1] Flipping bits in memory without accessing them
- [2] Google Project Zero: 1st Rowhammer root Exploit (flipping PTEs)
- [3] Rowhammer.js: 1st RH bit flip in JavaScript
- [4] Dedup est Machina: Breaking Microsoft Edge's sandbox
- [5] Flip Feng Shui: Breaking the cloud
- [6] Drammer: Flip feng shui goes mobile

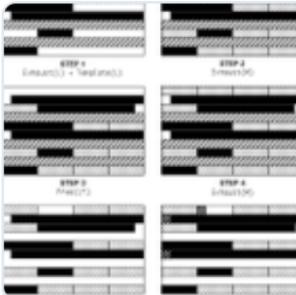
everyween
+100XP

GLitch: the chronicles



- [1] Flipping bits in memory without accessing them
- [2] Google Project Zero: 1st Rowhammer root Exploit (flipping PTEs)
- [3] Rowhammer.js: 1st RH bit flip in JavaScript
- [4] Dedup est Machina: Breaking Microsoft Edge's sandbox
- [5] Flip Feng Shui: Breaking the cloud
- [6] Drammer: Flip feng shui goes mobile

everyween
+100XP



Drammer: Flip Feng Shui Goes Mobile - VUsec

Drammer is the first instance of mobile Rowhammer and comprehends a deterministic Android root exploit that does not rely on any software vulnerability.

vusec.net

2

24

22

✉



Victor van der Veen @vvdveen · 25 ott 2016

I wouldn't be surprised if we could pull this one from a browser actually...

🌐 Traduci dalla lingua originale: inglese

1

4

5

✉



the grugq

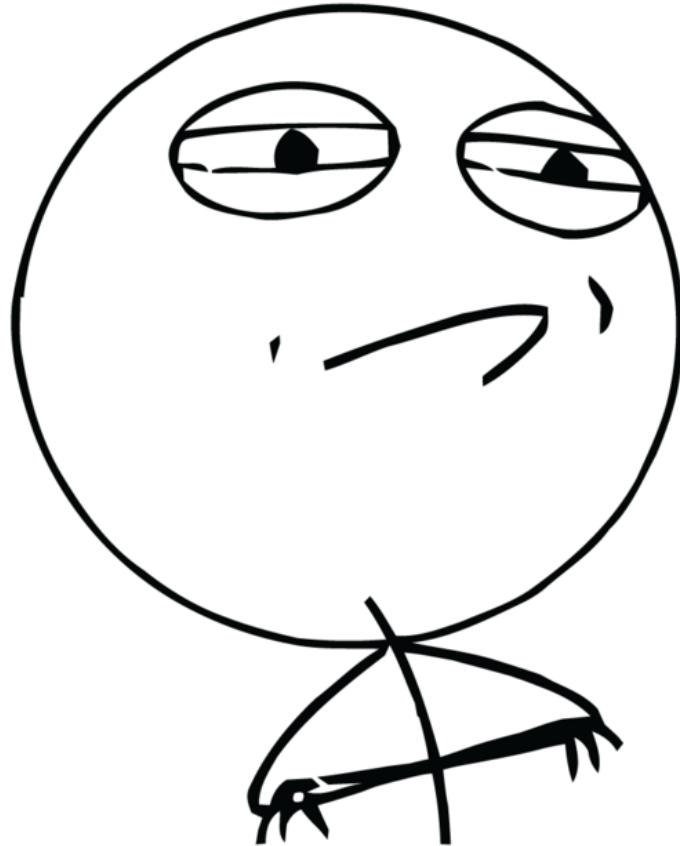
@thegrugq

Following

In risposta a @vvdveen e @vu5ec

love to see it happen. :)

CHALLENGE ACCEPTED



Attacker primitives

Attacker primitives

#P1. Fast cache eviction

- clflush (native)
- eviction sets (JS)

#P2. Contiguous memory

- THPs (native and JS)

Attacker primitives

#P1. Fast cache eviction

- clflush (native)
- eviction sets (JS)

#P2. Contiguous memory

- THPs (native and JS)



Attacker primitives

#P1. Fast cache eviction

✗ clflush (native)

✗ eviction sets (JS)

#P2. Contiguous memory

✗ THPs (native and JS)



Attacker primitives

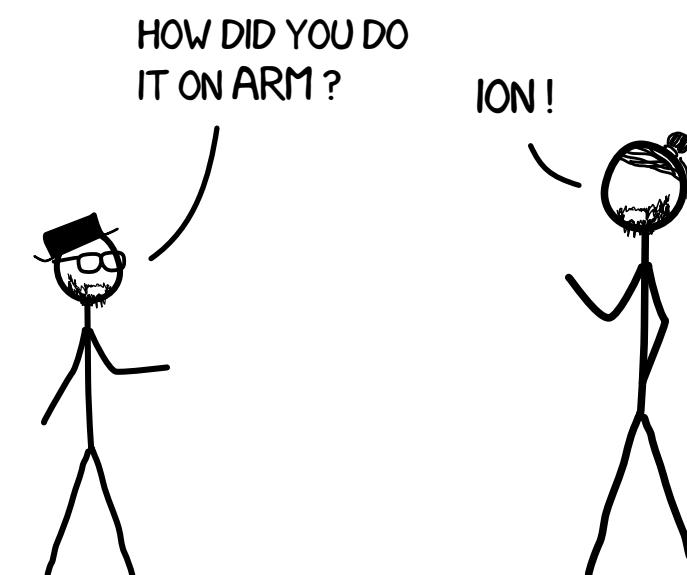
#P1. Fast cache eviction

✗ clflush (native)

✗ eviction sets (JS)

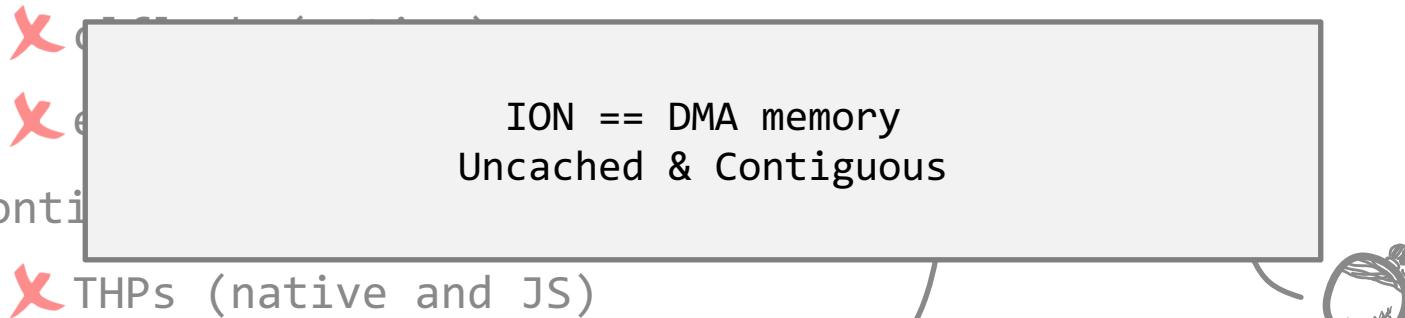
#P2. Contiguous memory

✗ THPs (native and JS)

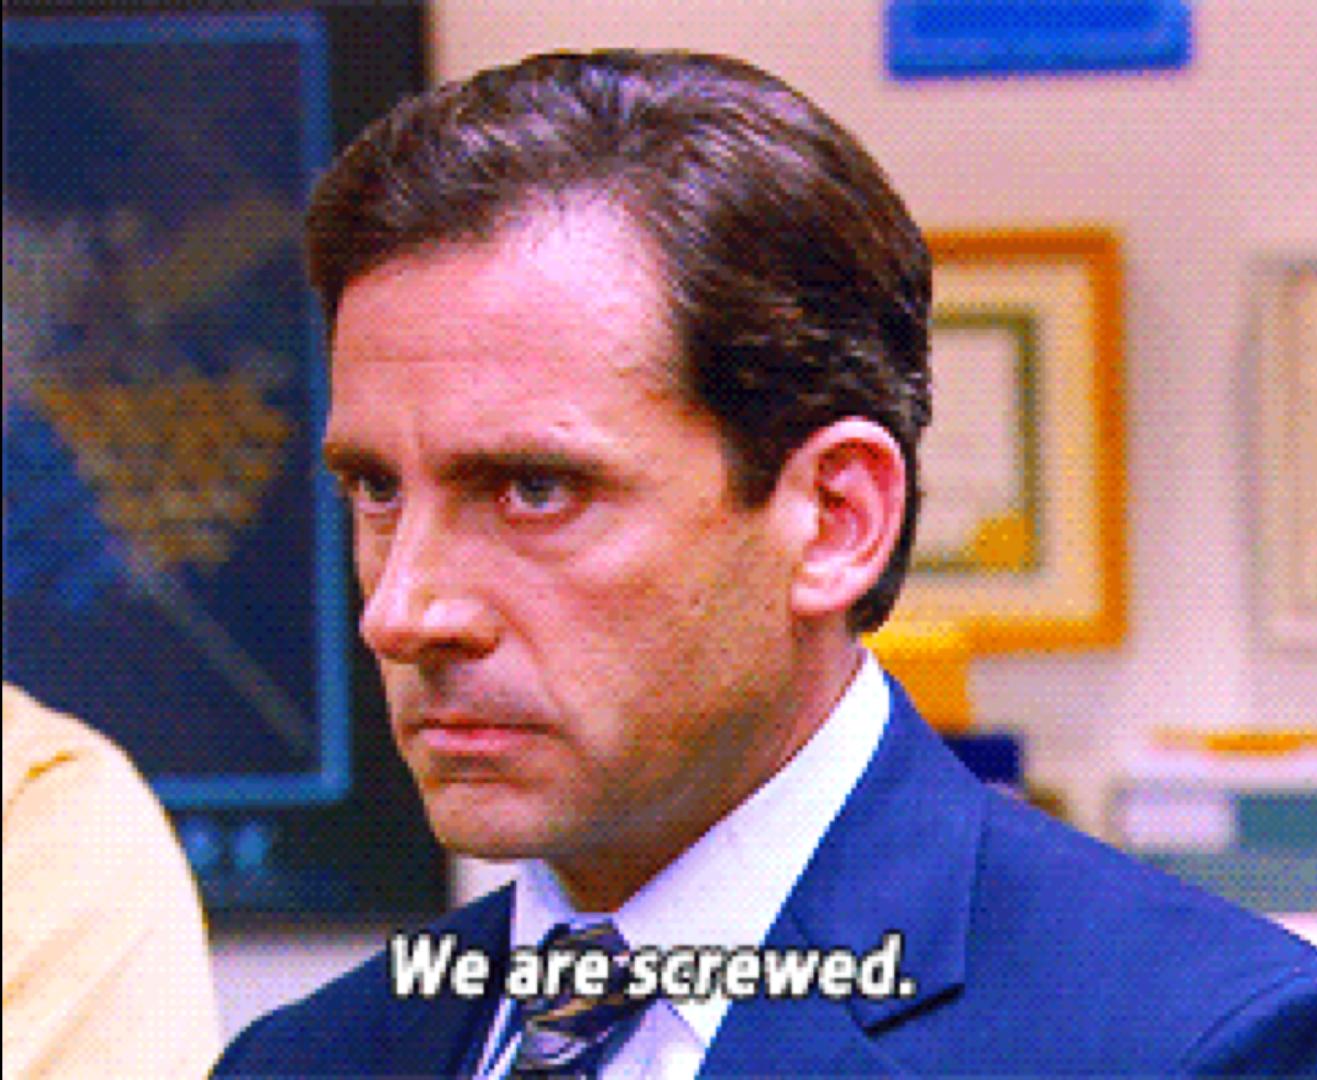


Attacker primitives

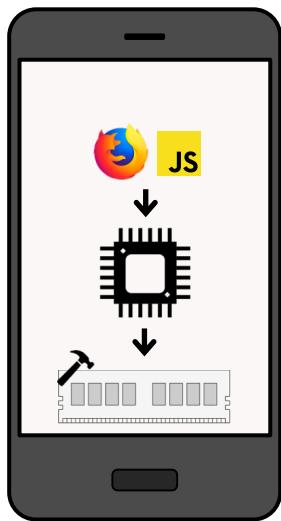
#P1. Fast cache eviction



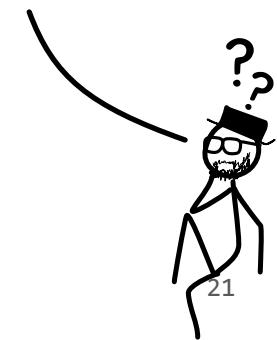
#P2. Contin



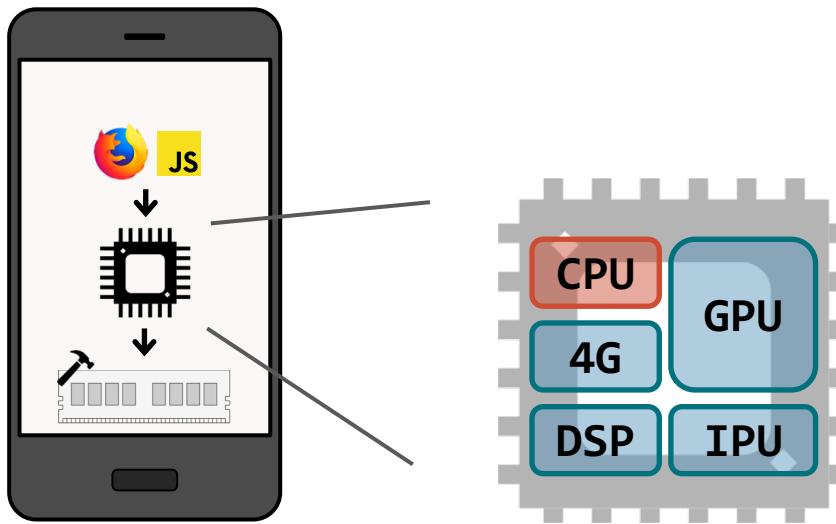
Attack Vector



WHAT IF YOU CHANGE
ATTACK VECTOR ?



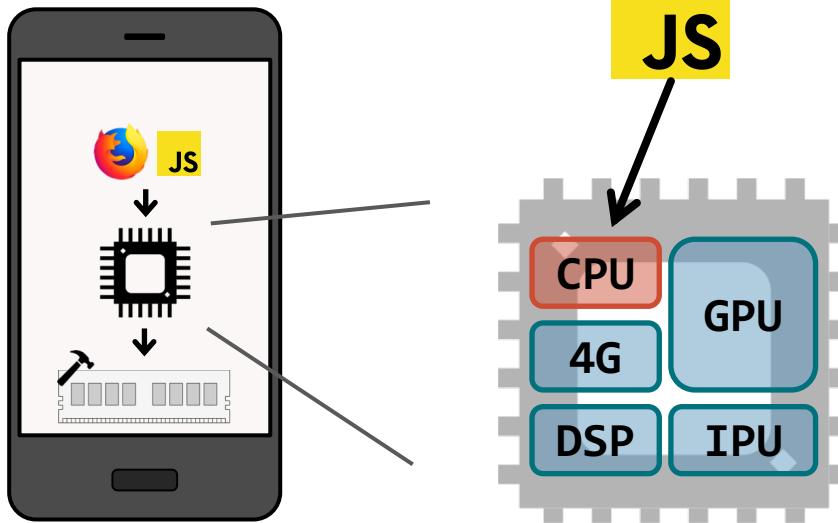
Attack Vector



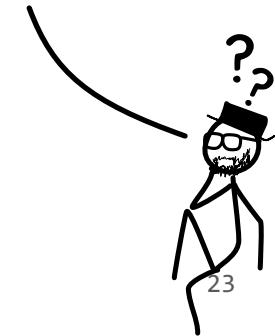
WHAT IF YOU CHANGE
ATTACK VECTOR ?



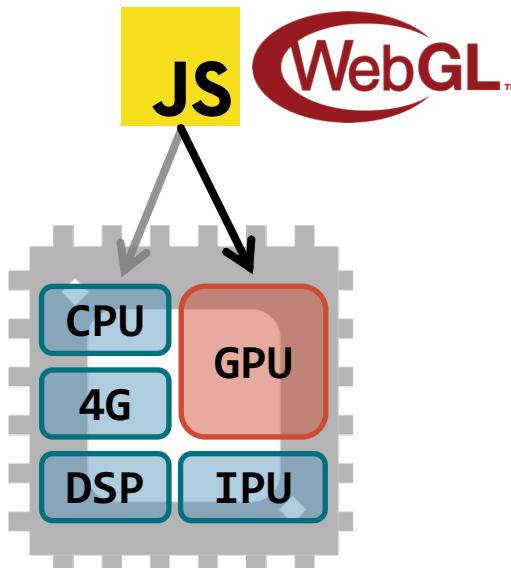
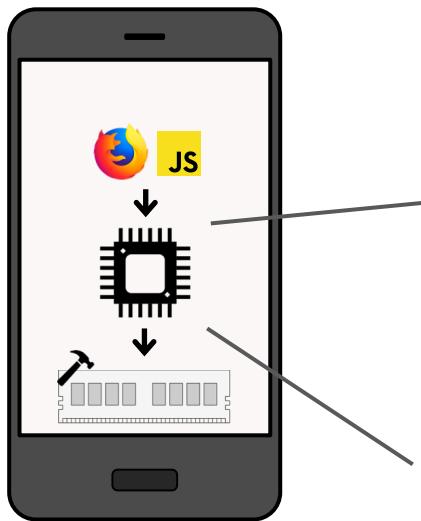
Attack Vector



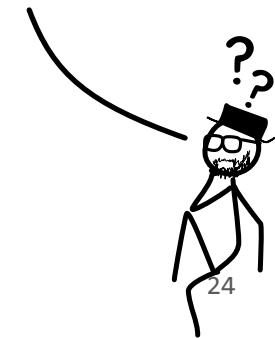
WHAT IF YOU CHANGE
ATTACK VECTOR ?



Attack Vector



WHAT IF YOU CHANGE
ATTACK VECTOR ?



Attacker primitives

#P1. DRAM Access

#P2. Fast cache eviction

#P3. Contiguous memory

Attacker primitives

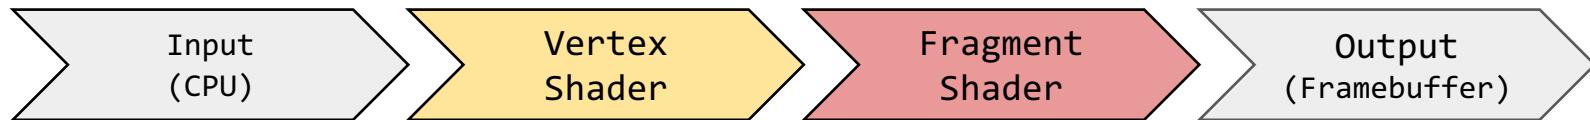
#P1. DRAM Access

#P2. Fast cache eviction

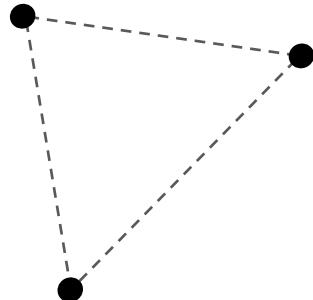
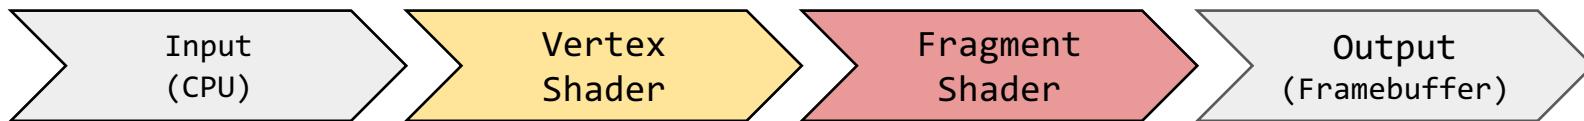
#P3. Contiguous memory

Understanding the GPU

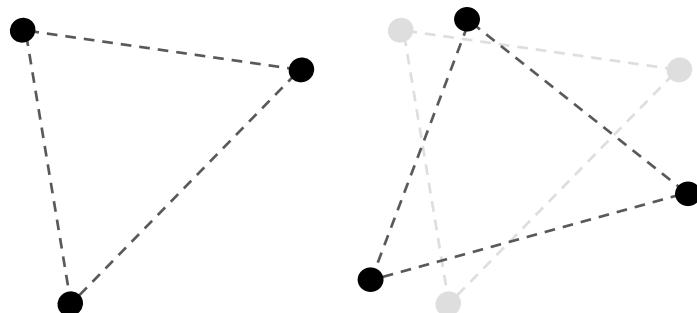
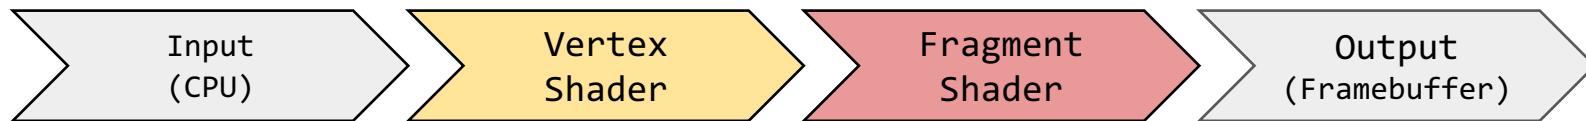
#P1. GPU: The rendering pipeline



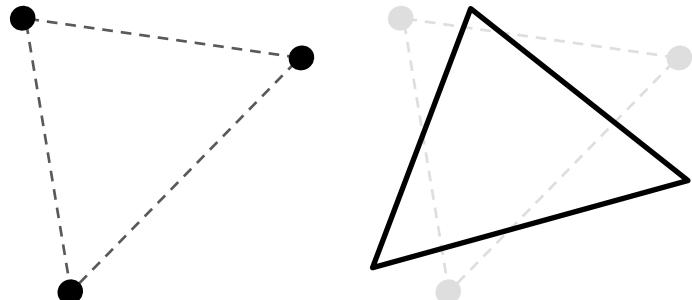
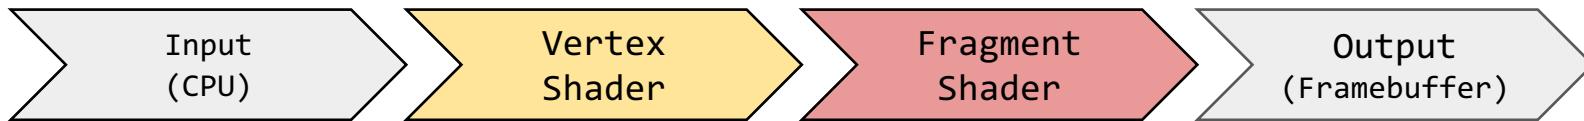
#P1. GPU: The rendering pipeline



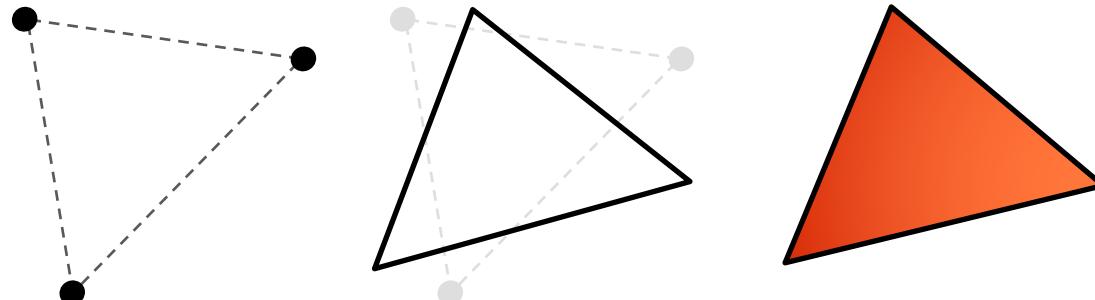
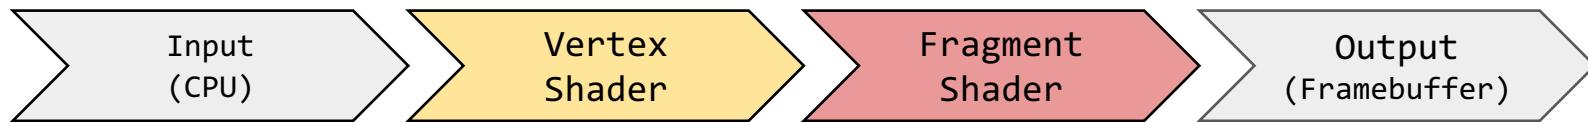
#P1. GPU: The rendering pipeline



#P1. GPU: The rendering pipeline



#P1. GPU: The rendering pipeline

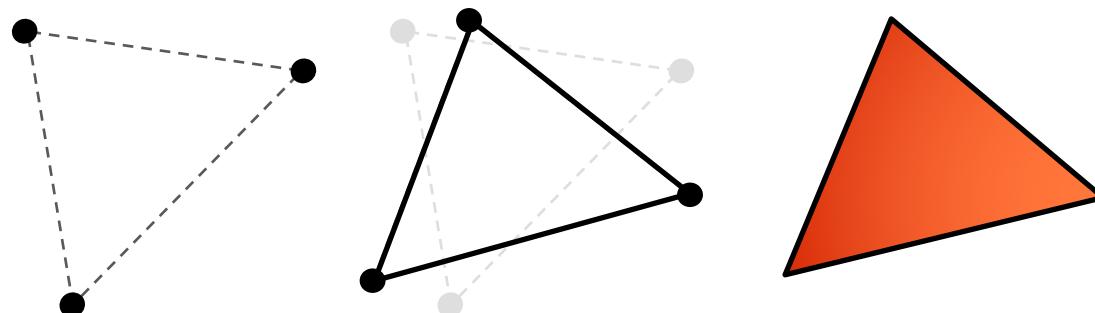
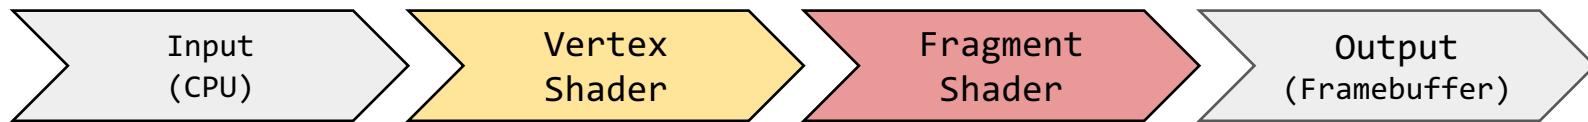


MINECRAFT

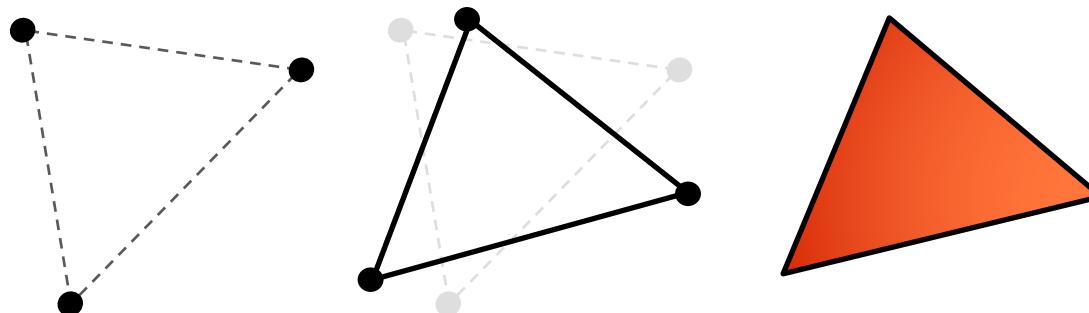
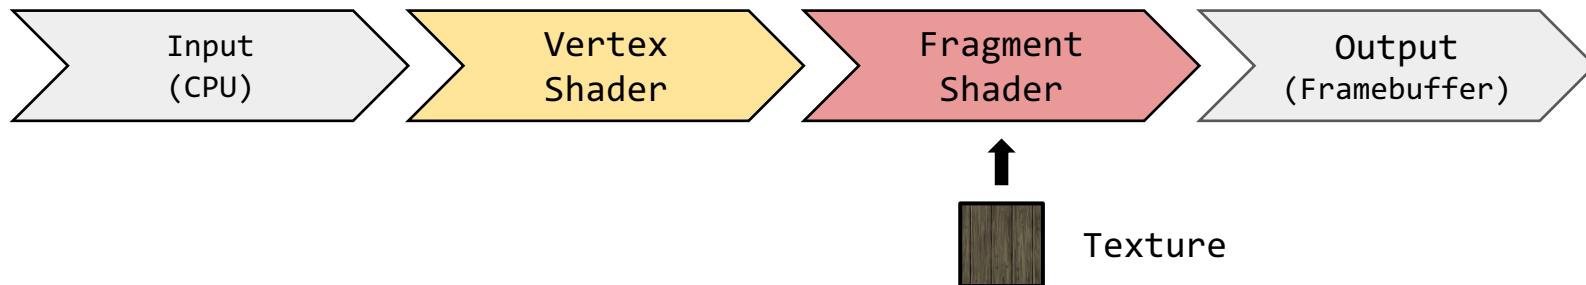
EDUCATION EDITION



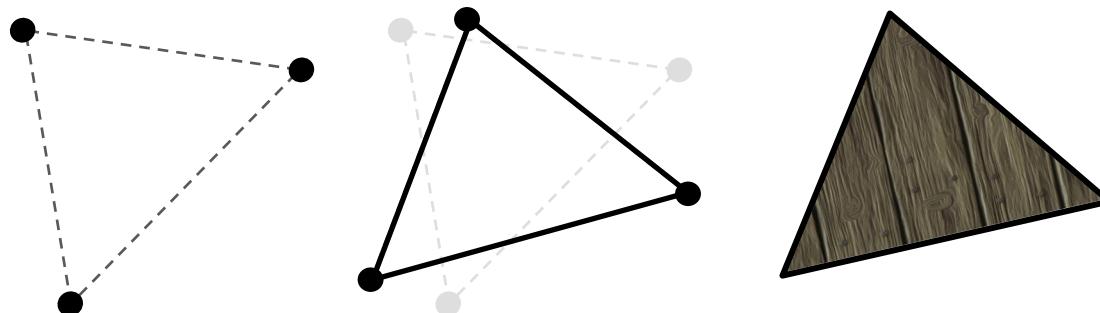
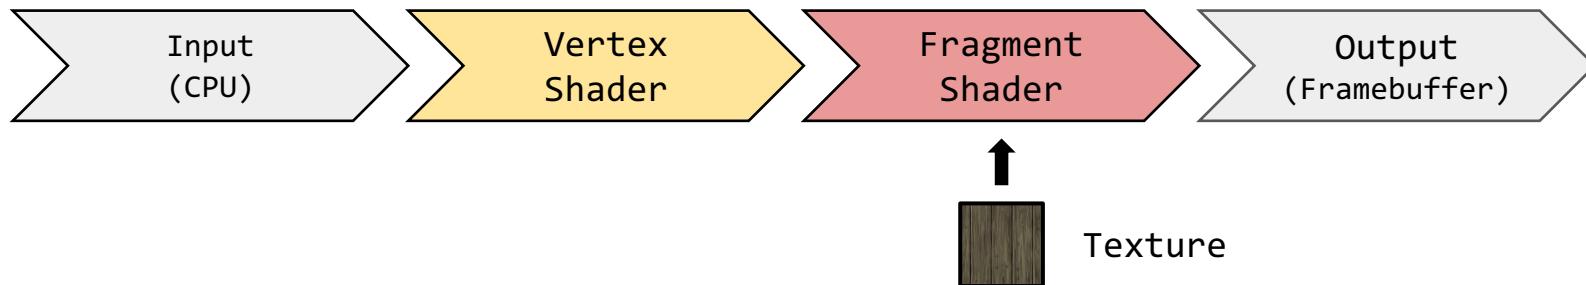
#P1. GPU: The rendering pipeline



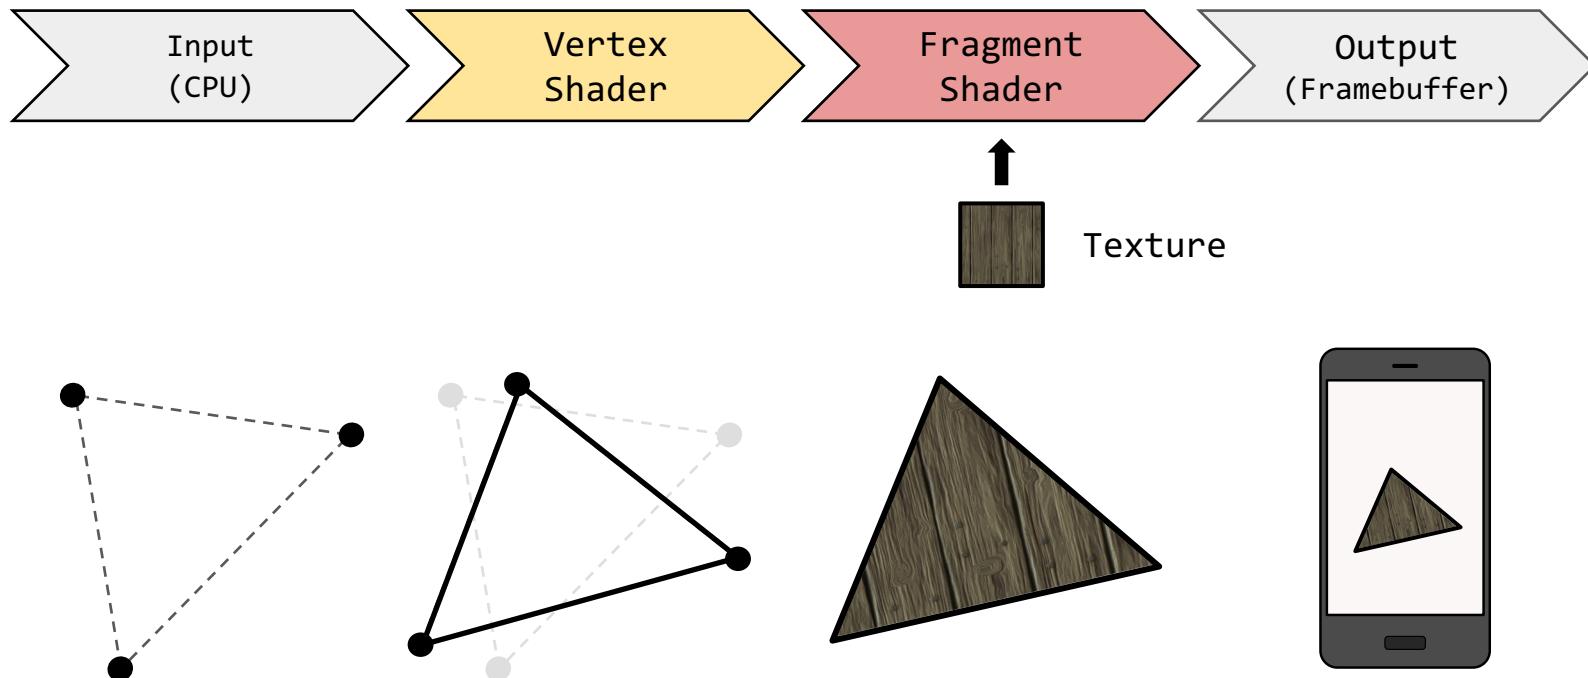
#P1. GPU: The rendering pipeline



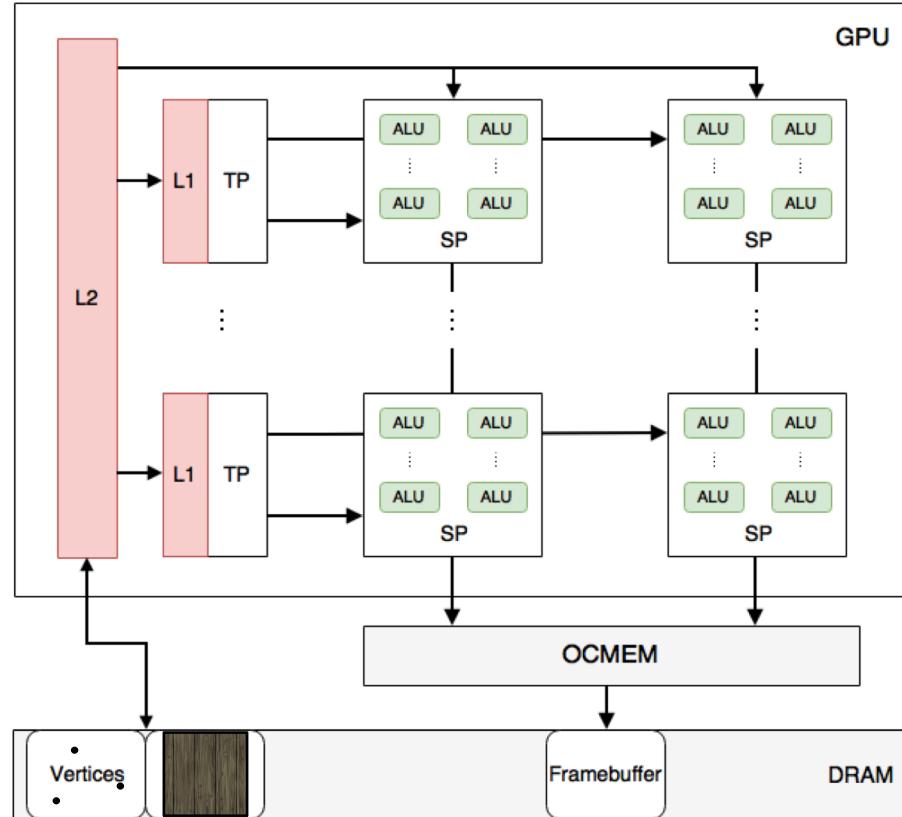
#P1. GPU: The rendering pipeline



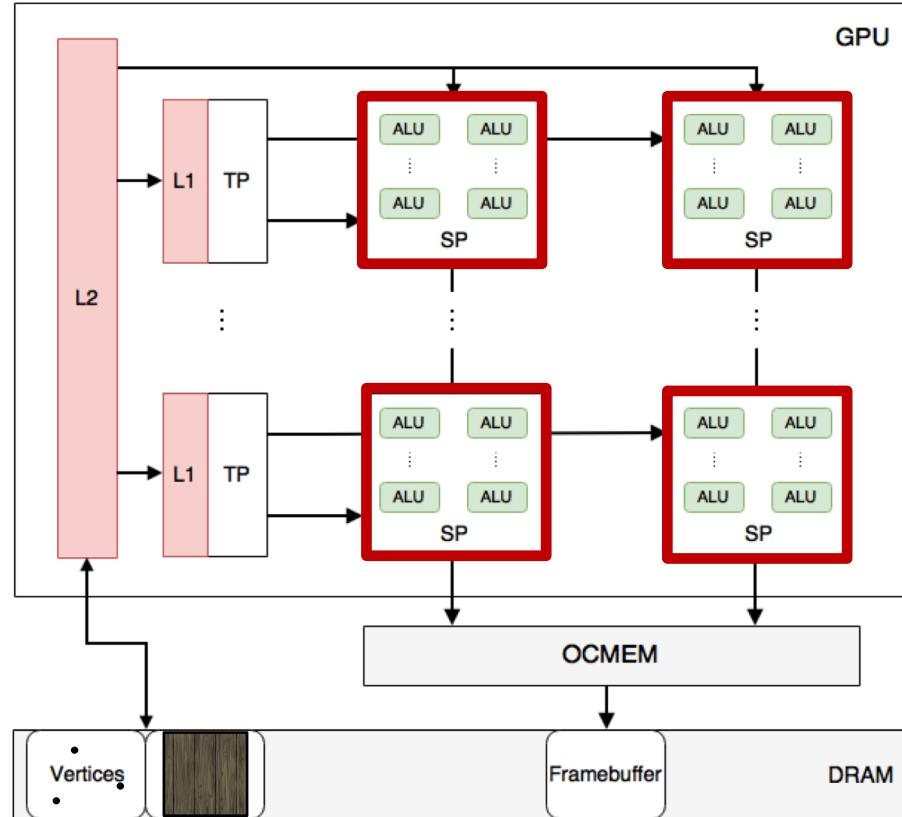
#P1. GPU: The rendering pipeline



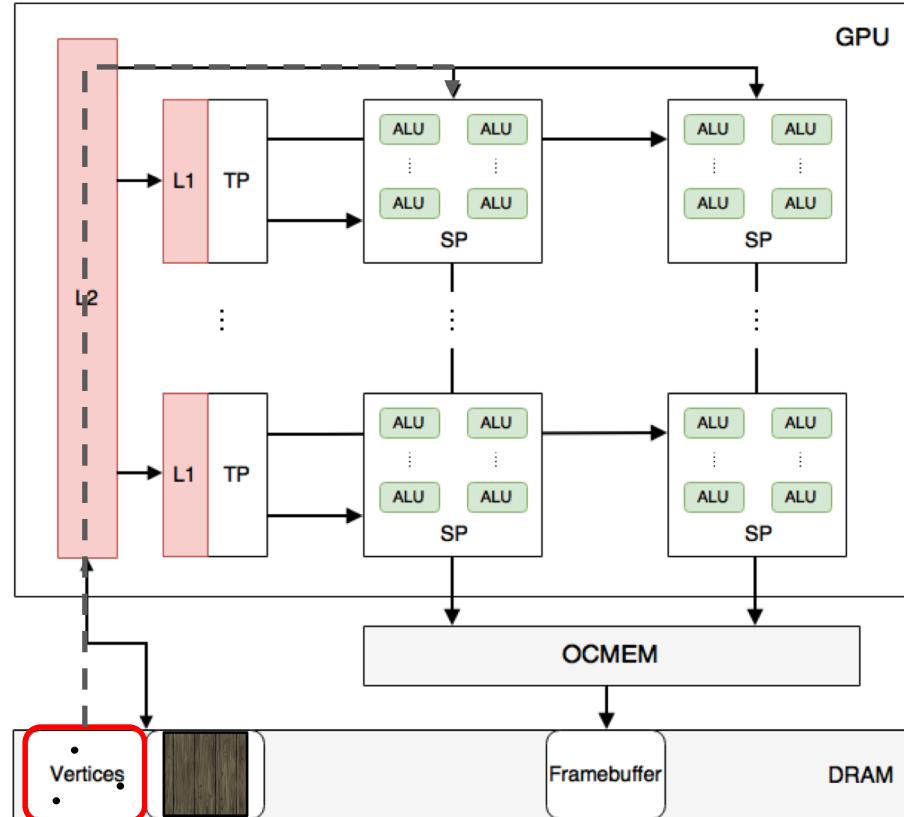
#P1. GPU: The architecture



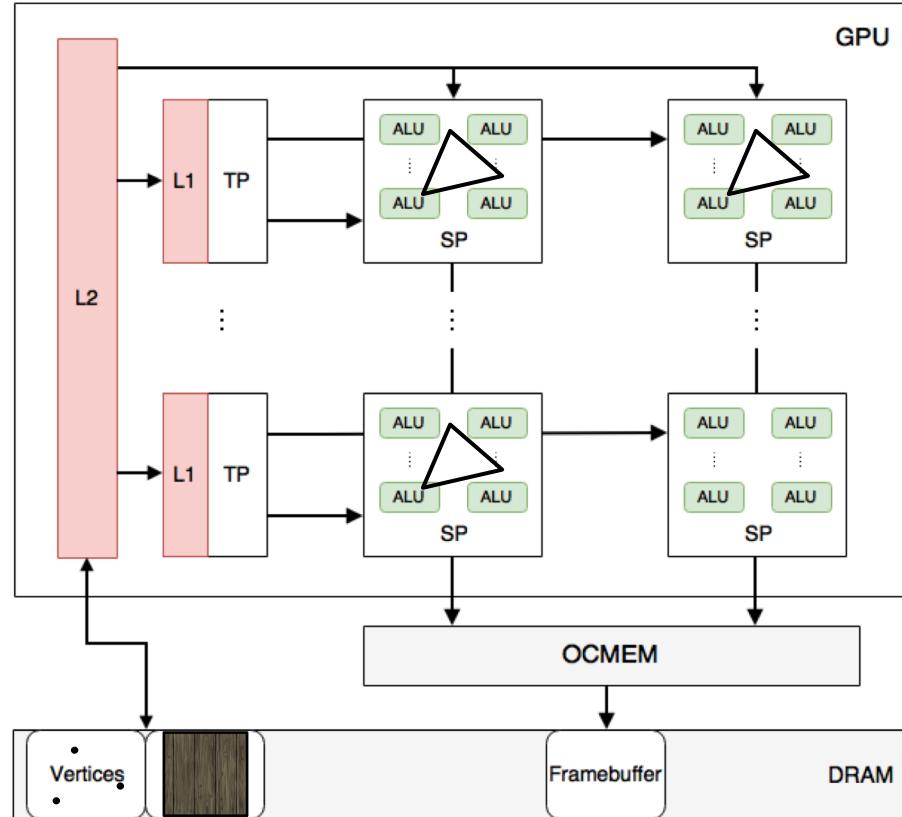
#P1. GPU: The architecture



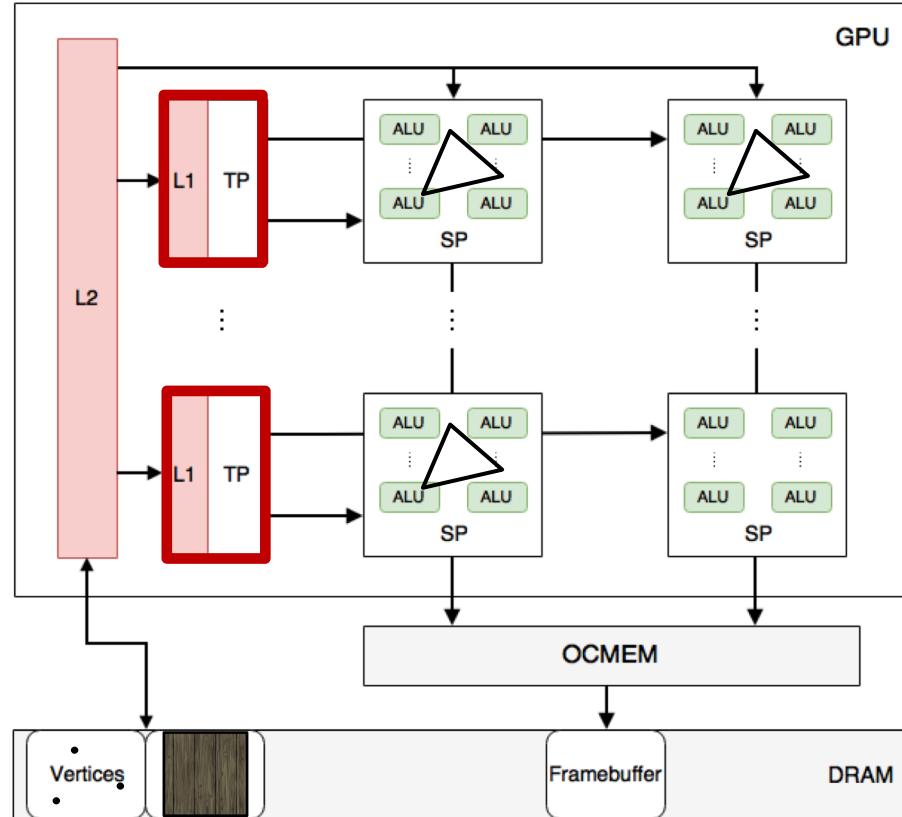
#P1. GPU: The architecture



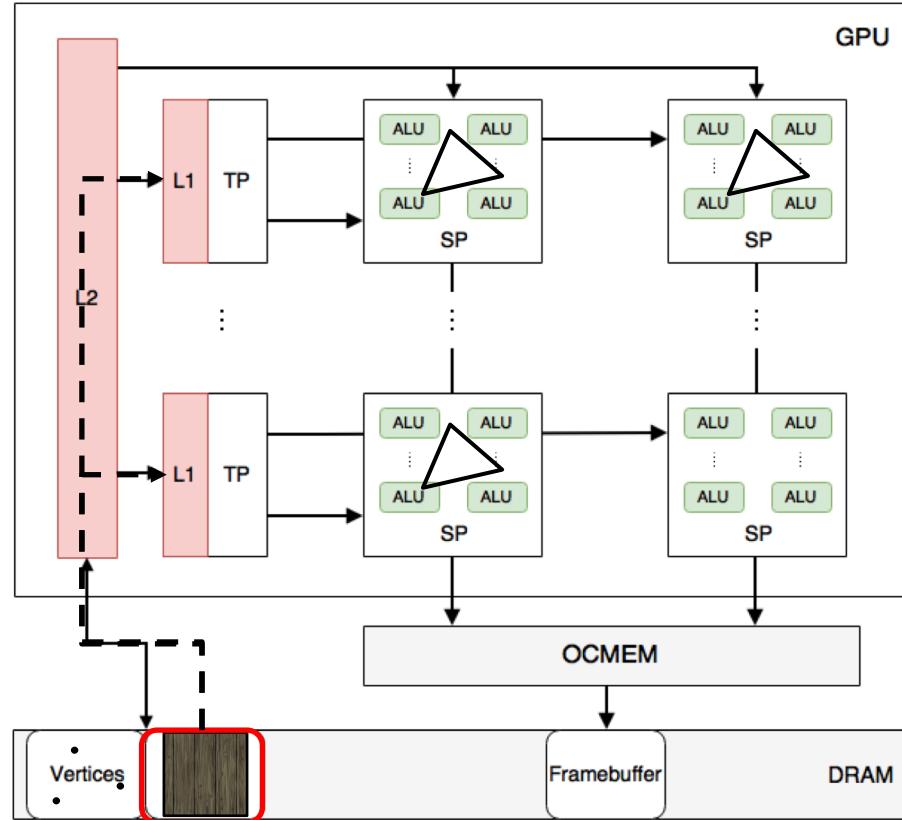
#P1. GPU: The architecture



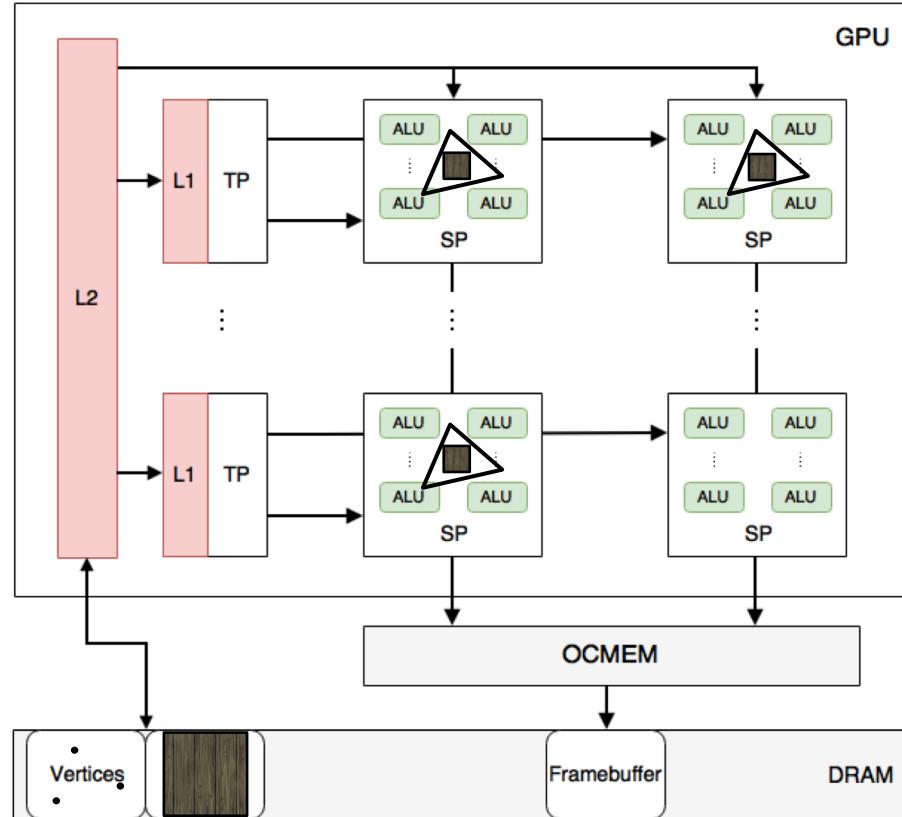
#P1. GPU: The architecture



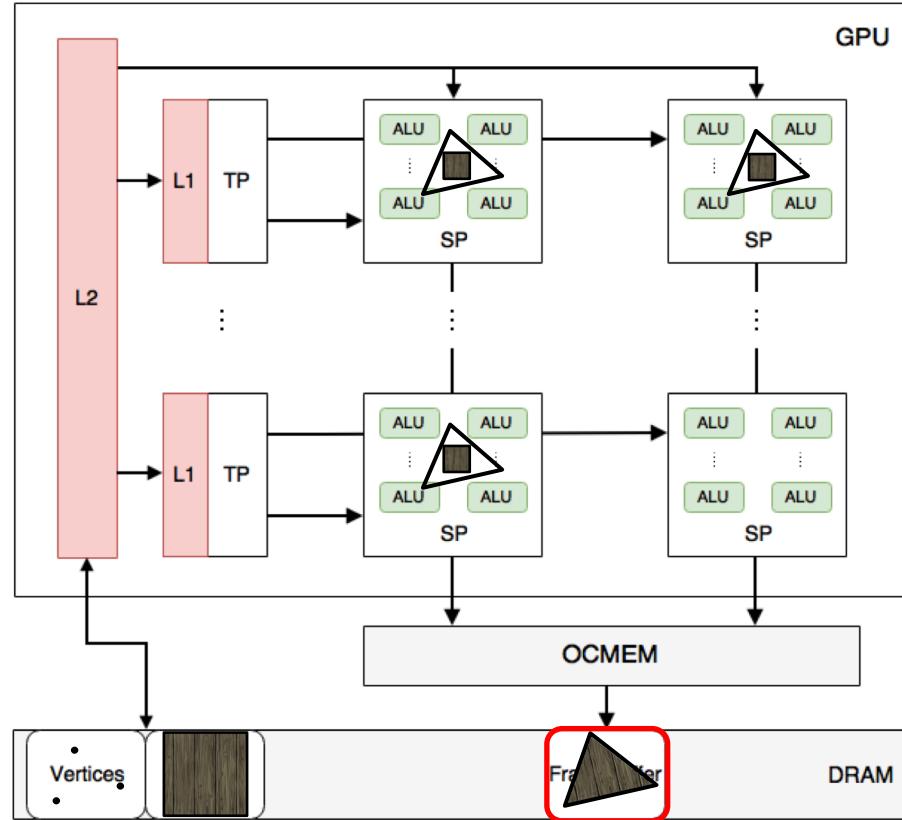
#P1. GPU: The architecture



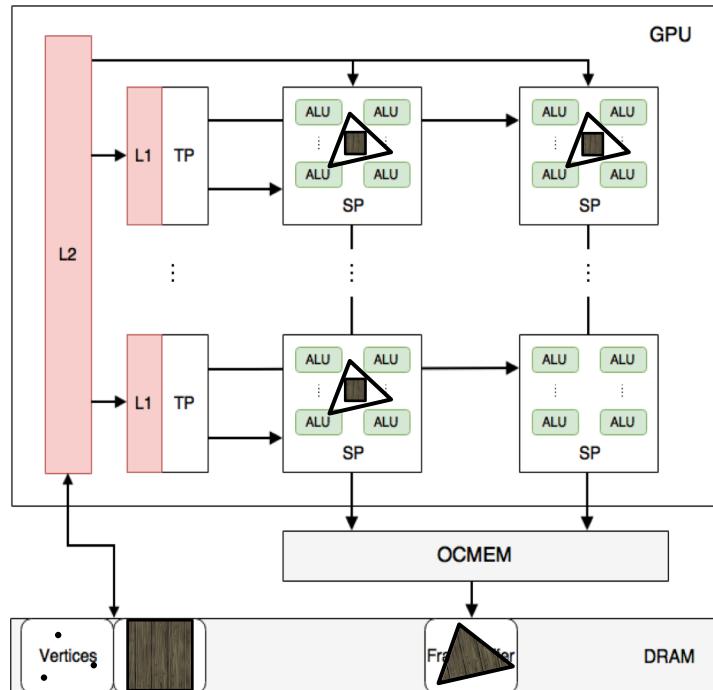
#P1. GPU: The architecture



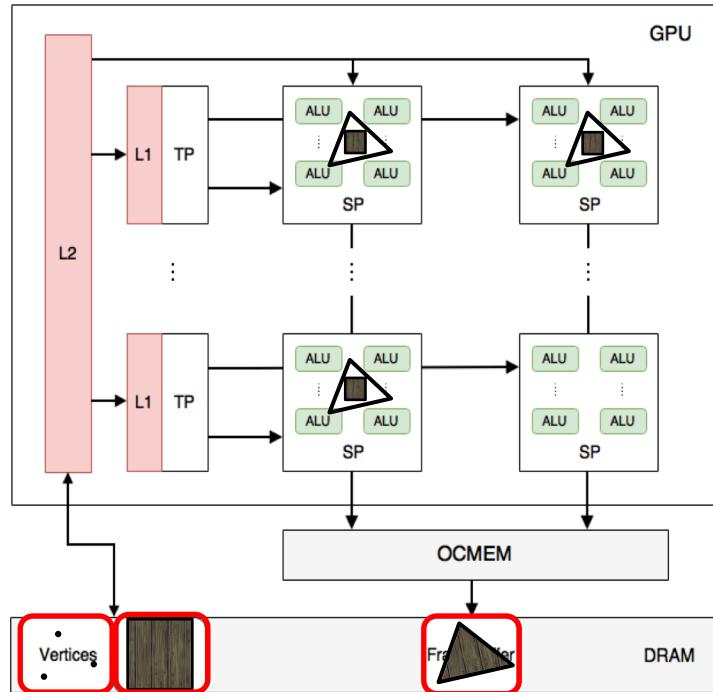
#P1. GPU: The architecture



#P1. DRAM access

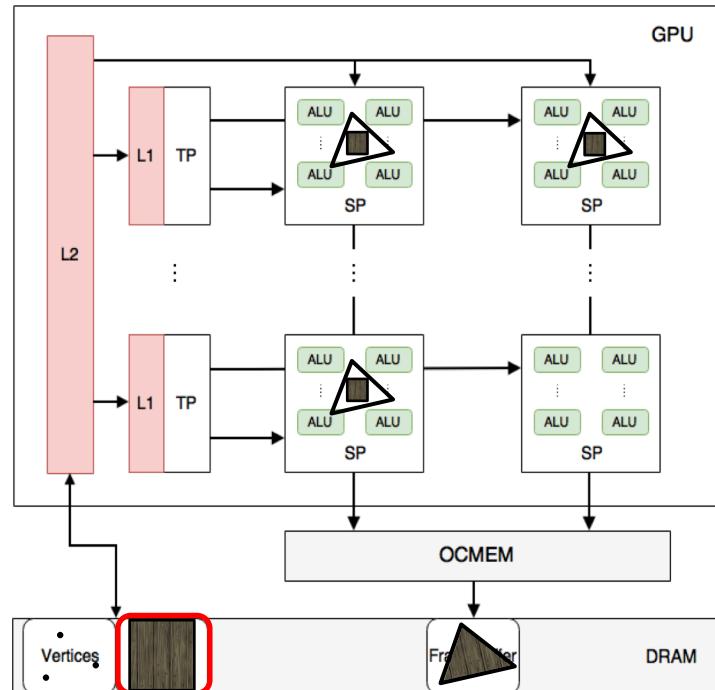


#P1. DRAM access



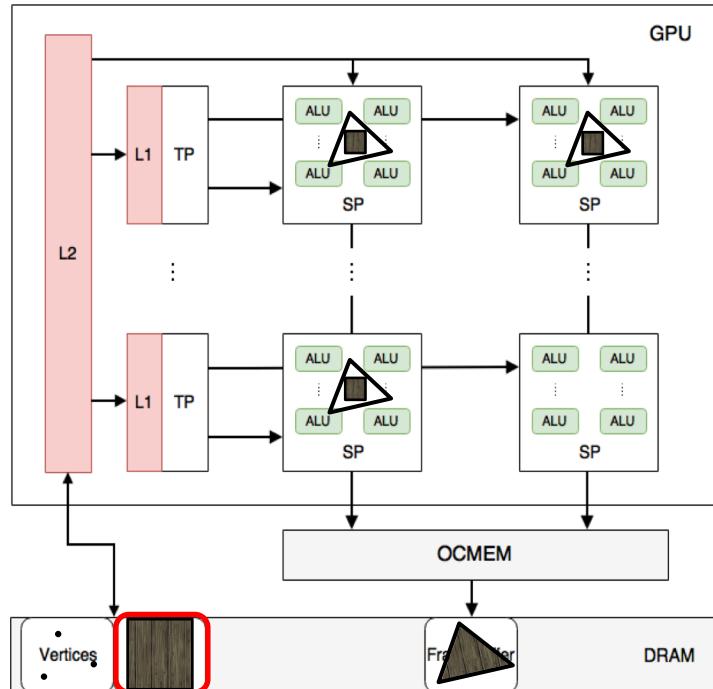
1. Read Vertices
2. Read Textures
3. Write to Framebuffer

#P1. DRAM access



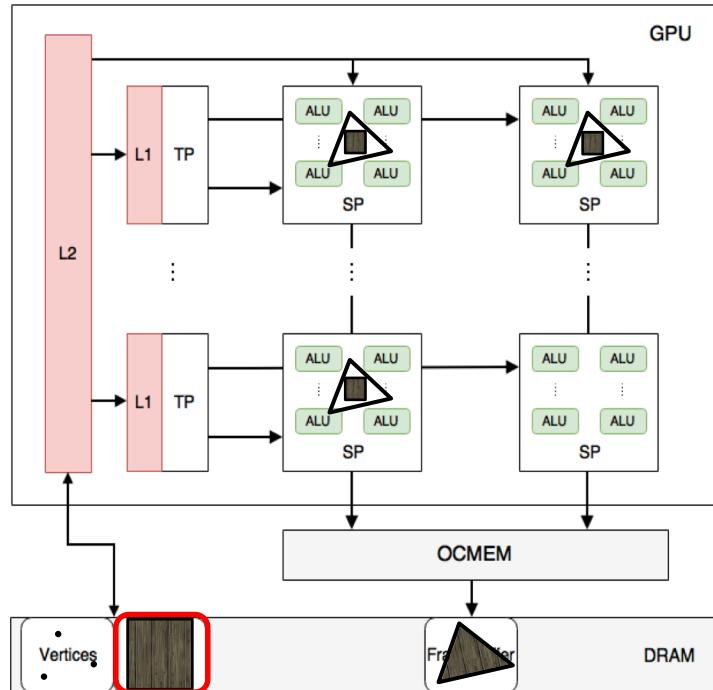
1. Read Vertices
2. Read Textures ==> most predictable
3. Write to Framebuffer

#P1. DRAM access: texture sampling



```
1 uniform sampler2D tex;  
2  
3 void main() {  
4     vec2 coord = vec2(0,0);  
5     gl_FragColor = texture2D(tex, coord);  
6 }
```

#P1. DRAM access: texture sampling



```
1 uniform sampler2D tex;  
2  
3 void main() {  
4     vec2 coord = vec2(0,0);  
5     gl_FragColor = texture2D(tex, coord);  
6 }  
                           tex[coord]
```

Attacker primitives

#P1. DRAM access ✓

#P2. Fast cache eviction

#P3. Contiguous memory

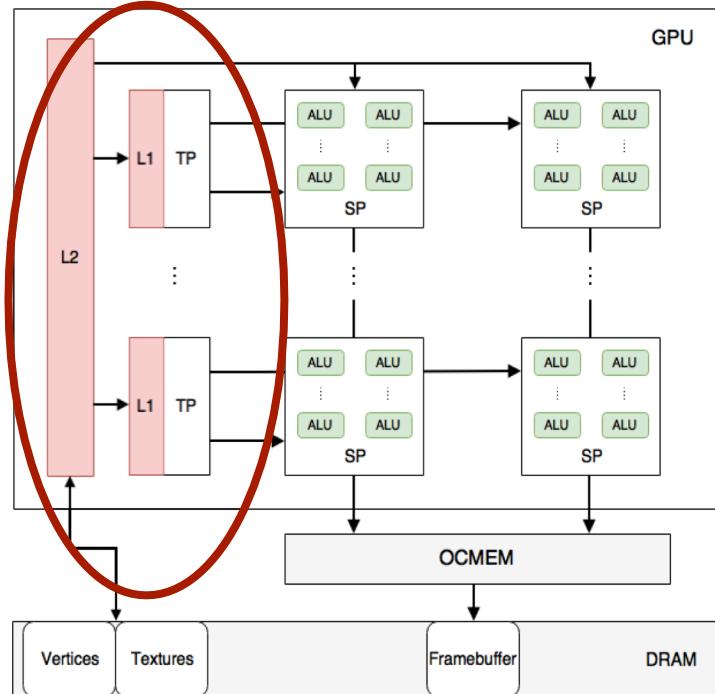
Attacker primitives

#P1. DRAM access ✓

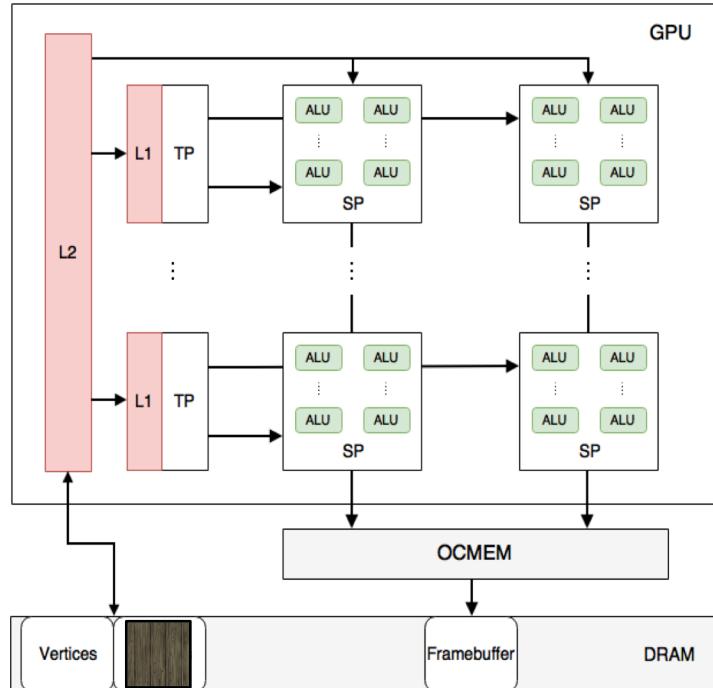
#P2. Fast cache eviction

#P3. Contiguous memory

#P2. Fast cache eviction

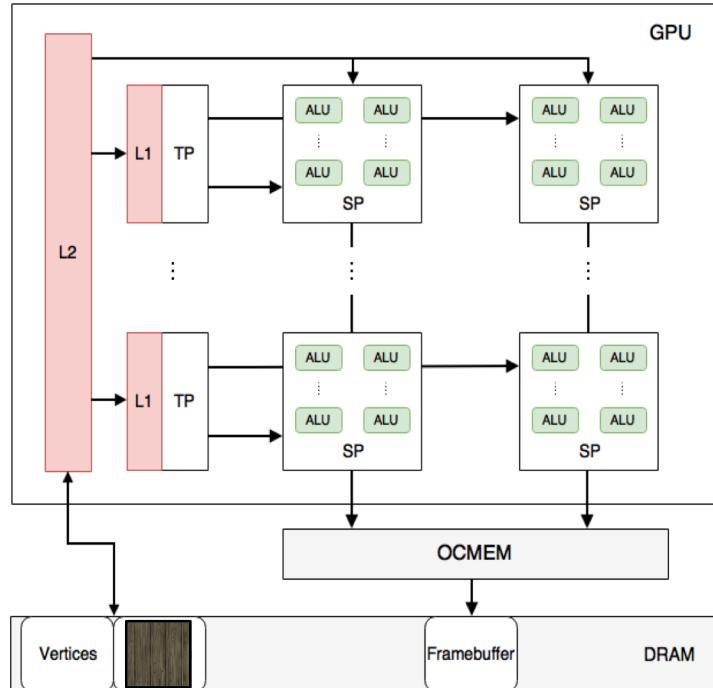


#P2. Fast cache eviction



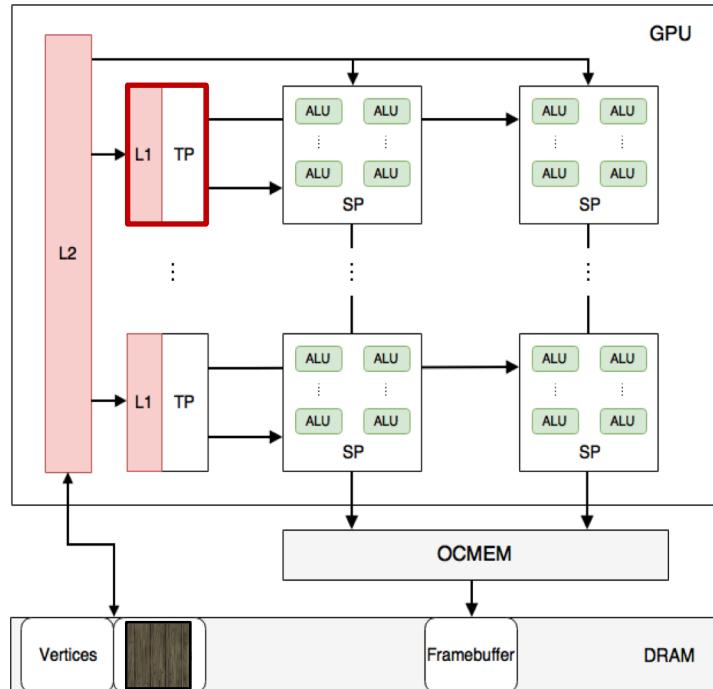
```
1 uniform sampler2D tex;  
2  
3 void main() {  
4     vec2 coord = vec2(0,0);  
5     vec4 a = texture2D(tex, coord);  
6     vec4 b = texture2D(tex, coord);  
7 }
```

#P2. Fast cache eviction



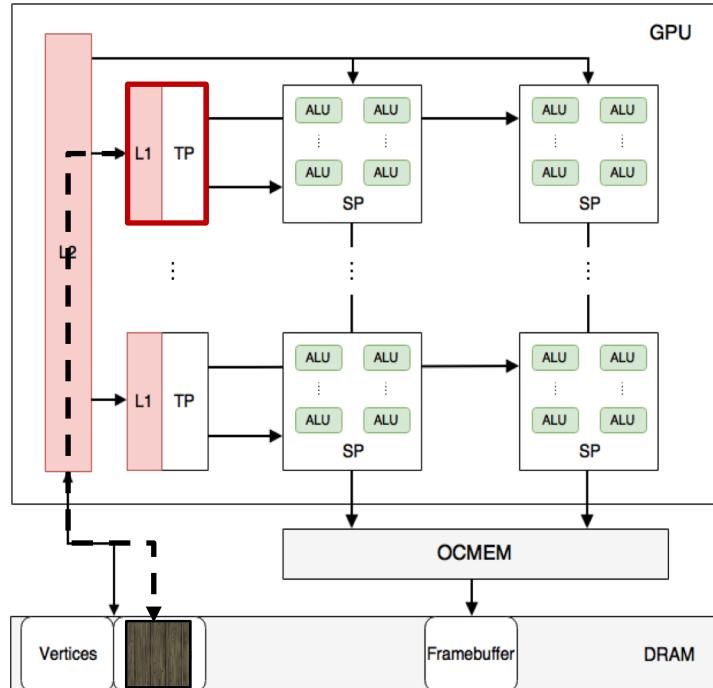
```
1 uniform sampler2D tex;  
2  
3 void main() {  
4     vec2 coord = vec2(0,0);  
5     vec4 a = texture2D(tex, coord); ←  
6     vec4 b = texture2D(tex, coord);  
7 }
```

#P2. Fast cache eviction



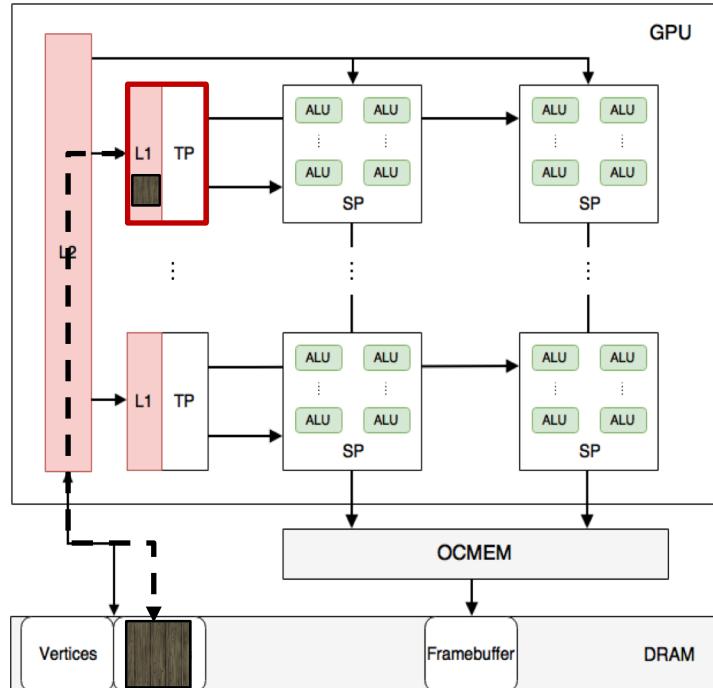
```
1 uniform sampler2D tex;  
2  
3 void main() {  
4     vec2 coord = vec2(0,0);  
5     vec4 a = texture2D(tex, coord); ←  
6     vec4 b = texture2D(tex, coord);  
7 }
```

#P2. Fast cache eviction



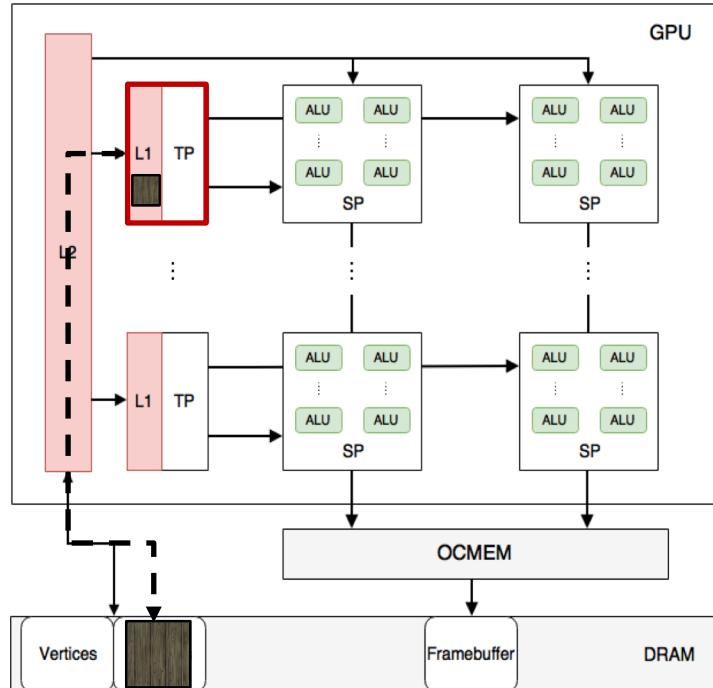
```
1 uniform sampler2D tex;
2
3 void main() {
4     vec2 coord = vec2(0,0);
5     vec4 a = texture2D(tex, coord); ←
6     vec4 b = texture2D(tex, coord);
7 }
```

#P2. Fast cache eviction



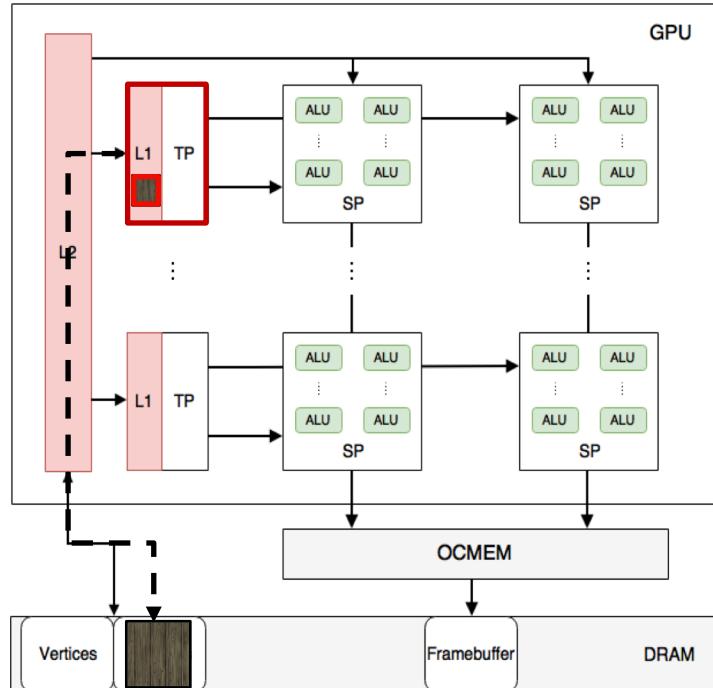
```
1 uniform sampler2D tex;
2
3 void main() {
4     vec2 coord = vec2(0,0);
5     vec4 a = texture2D(tex, coord); ←
6     vec4 b = texture2D(tex, coord);
7 }
```

#P2. Fast cache eviction



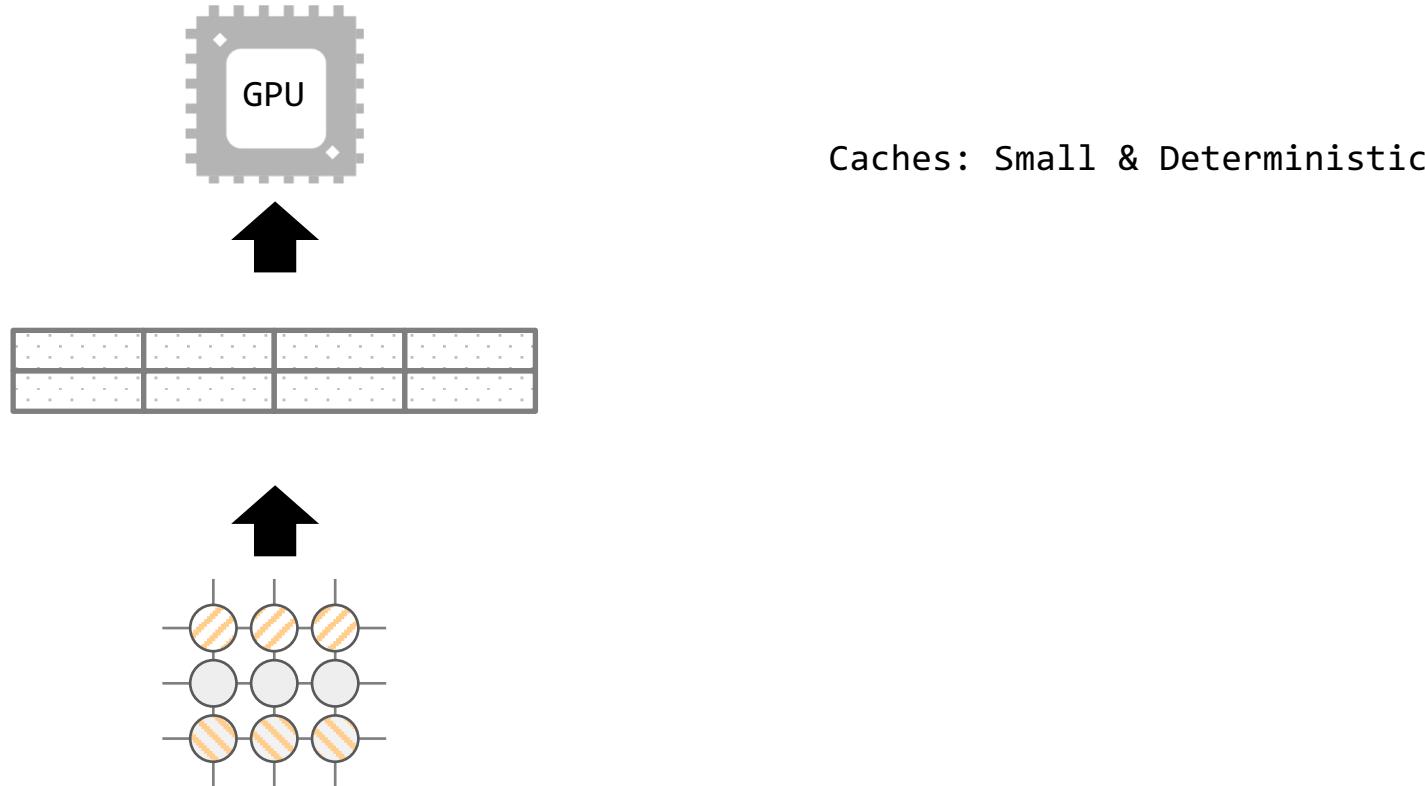
```
1 uniform sampler2D tex;  
2  
3 void main() {  
4     vec2 coord = vec2(0,0);  
5     vec4 a = texture2D(tex, coord);  
6     vec4 b = texture2D(tex, coord); ←  
7 }
```

#P2. Fast cache eviction

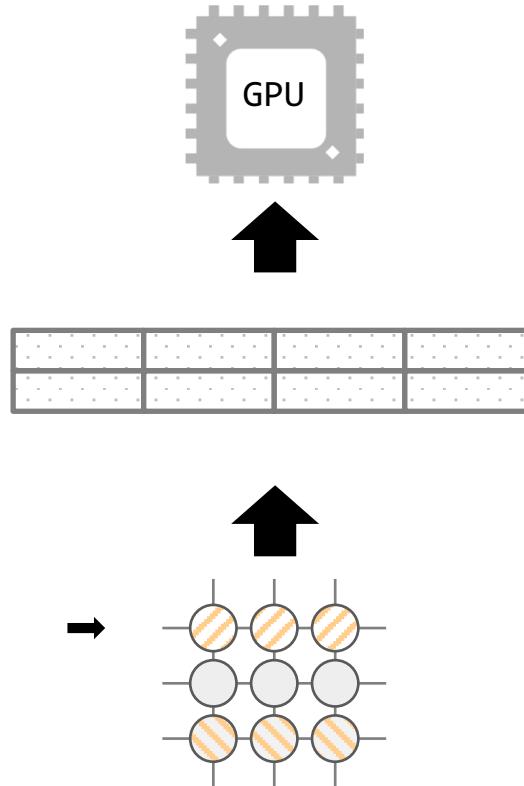


```
1 uniform sampler2D tex;  
2  
3 void main() {  
4     vec2 coord = vec2(0,0);  
5     vec4 a = texture2D(tex, coord);  
6     vec4 b = texture2D(tex, coord); ←  
7 }
```

#P2. Eviction-based Rowhammer: GPU



#P2. Eviction-based Rowhammer: GPU

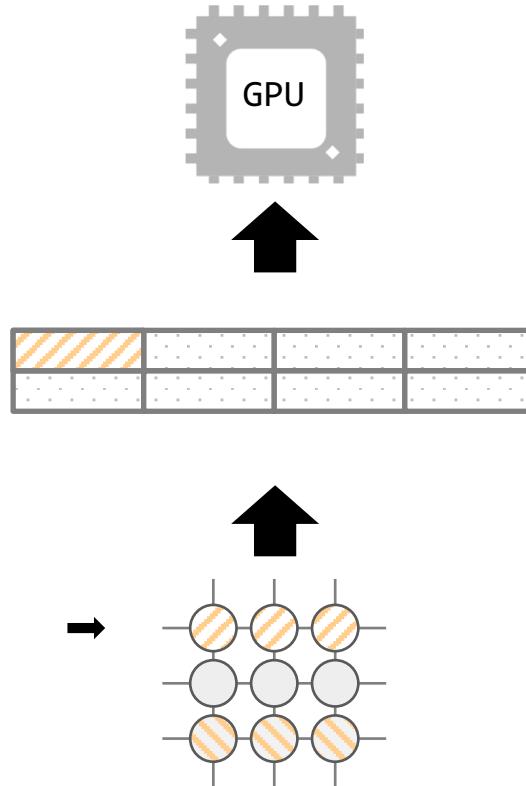


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU

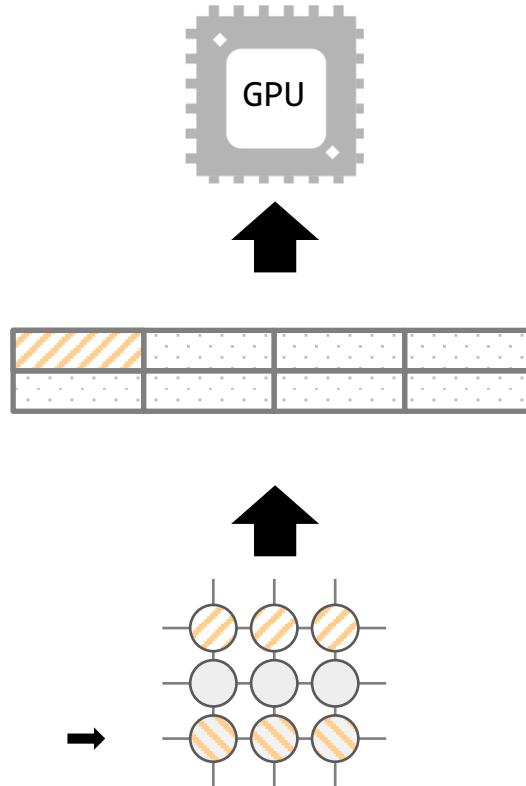


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU

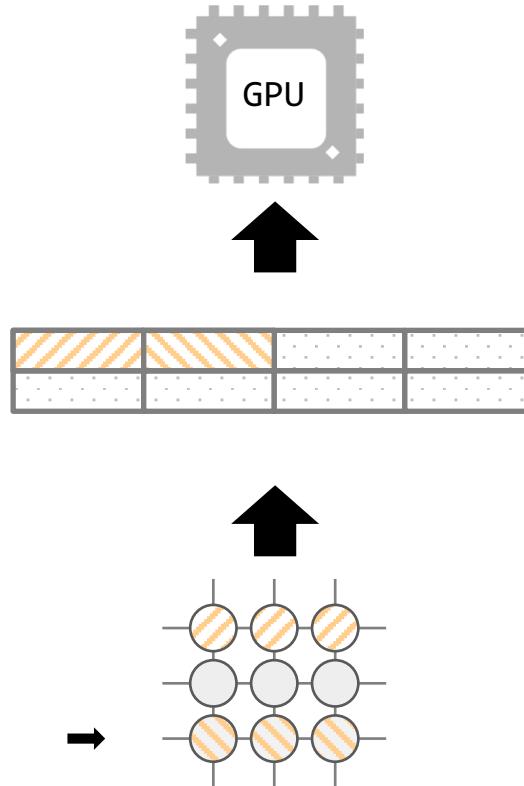


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU

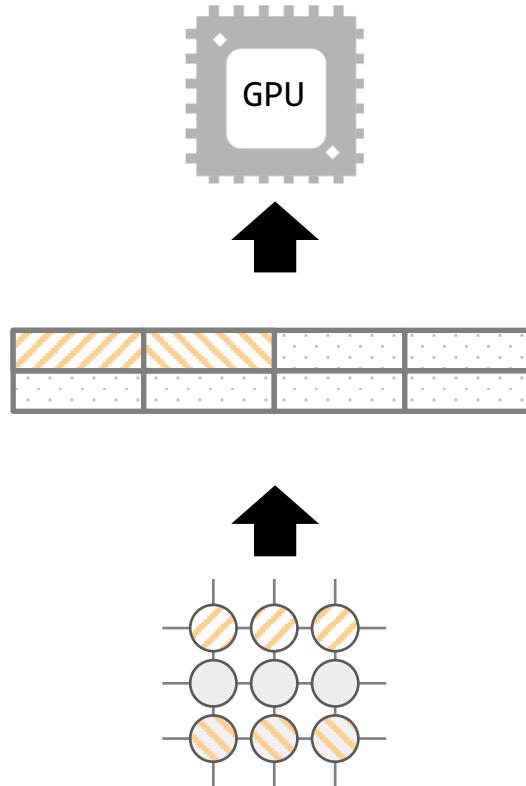


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU

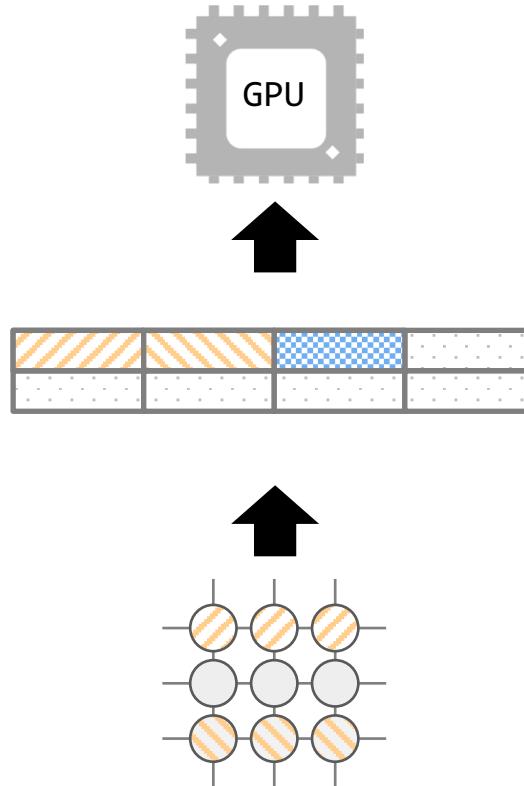


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU

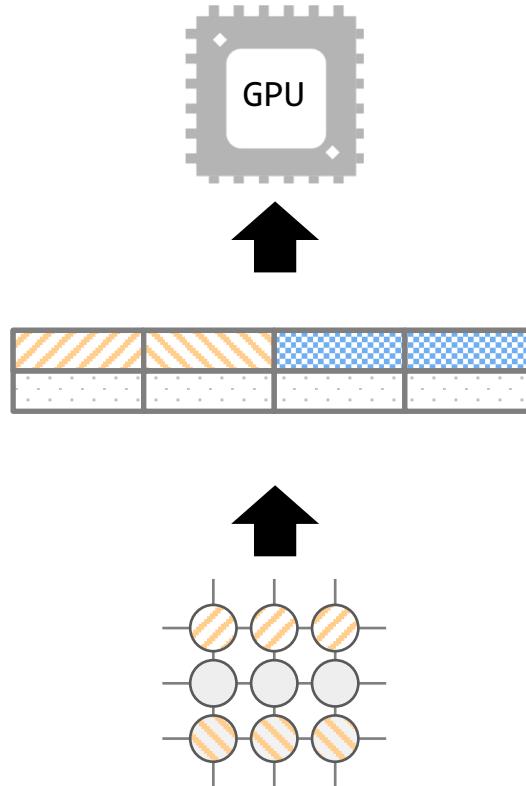


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU

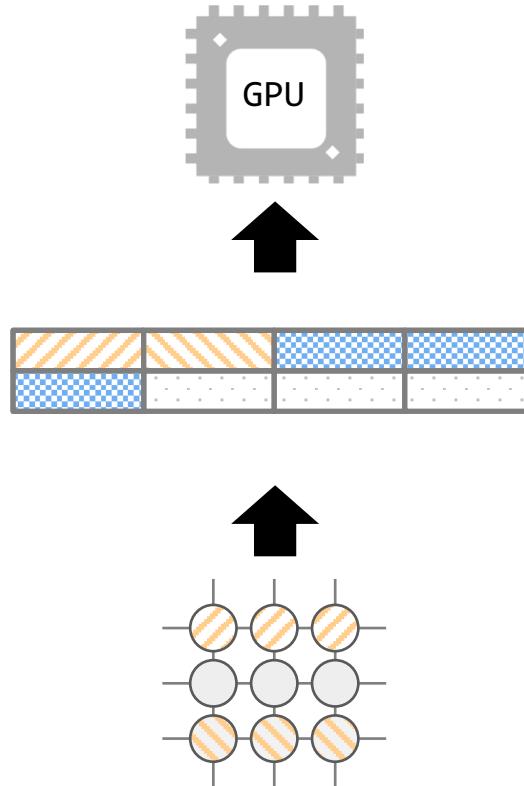


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU

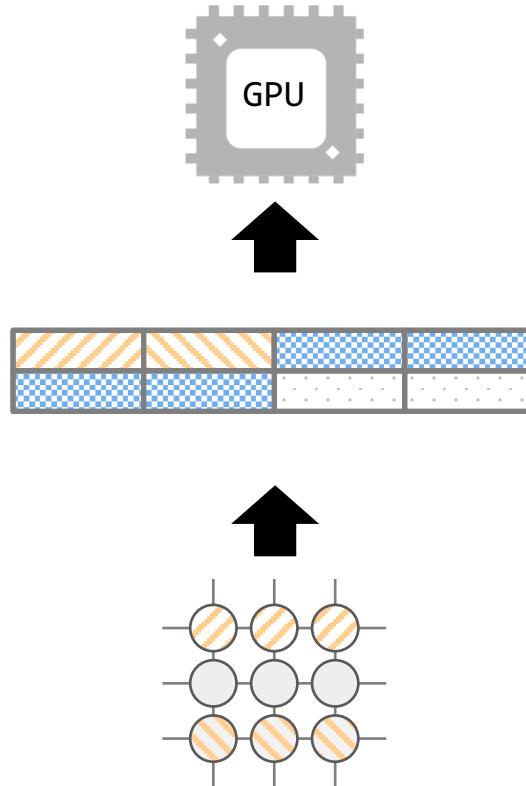


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU

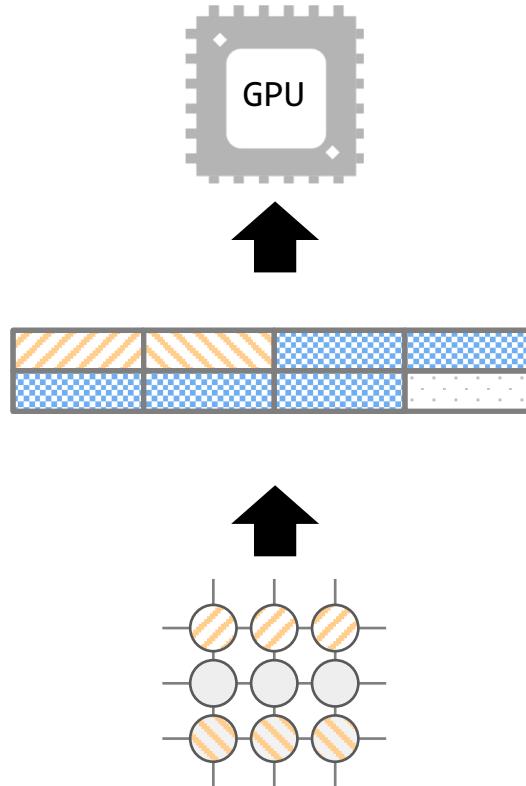


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU

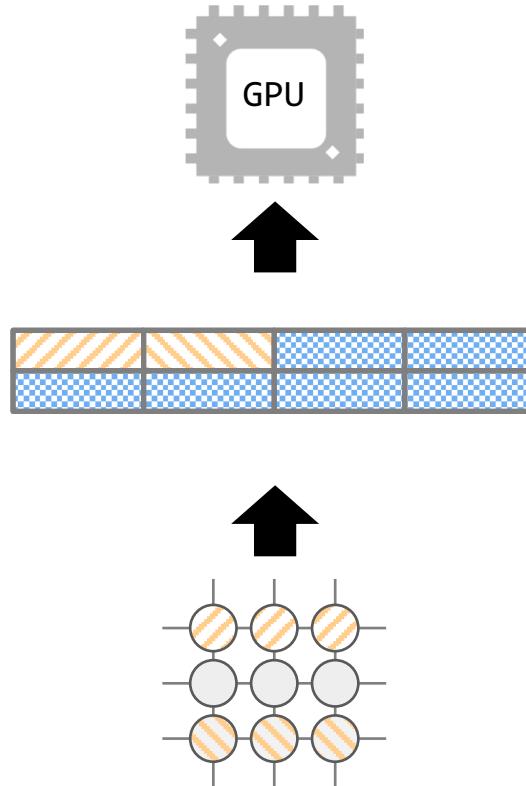


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU

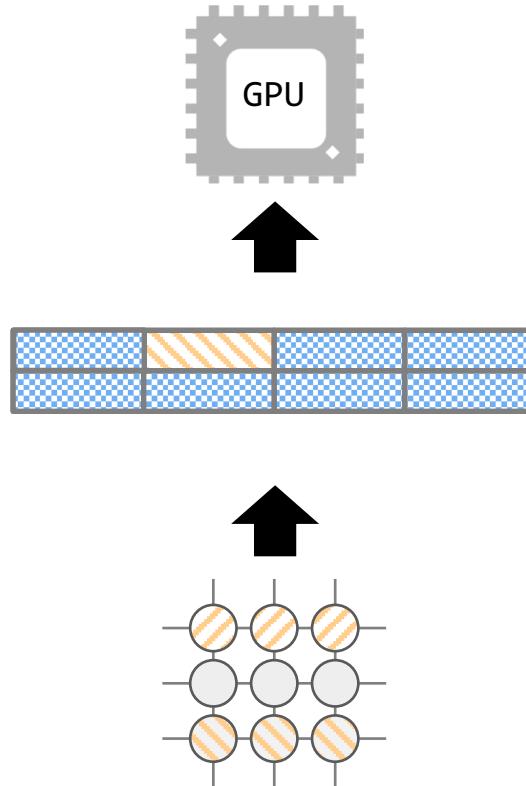


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU

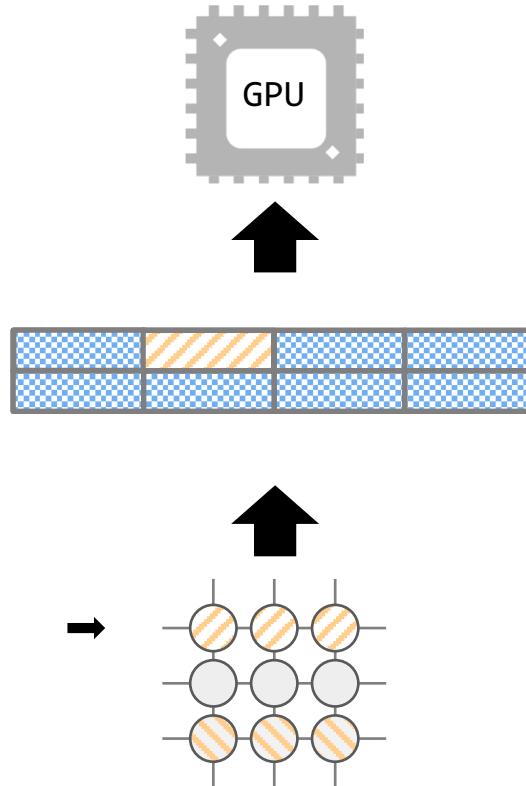


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU

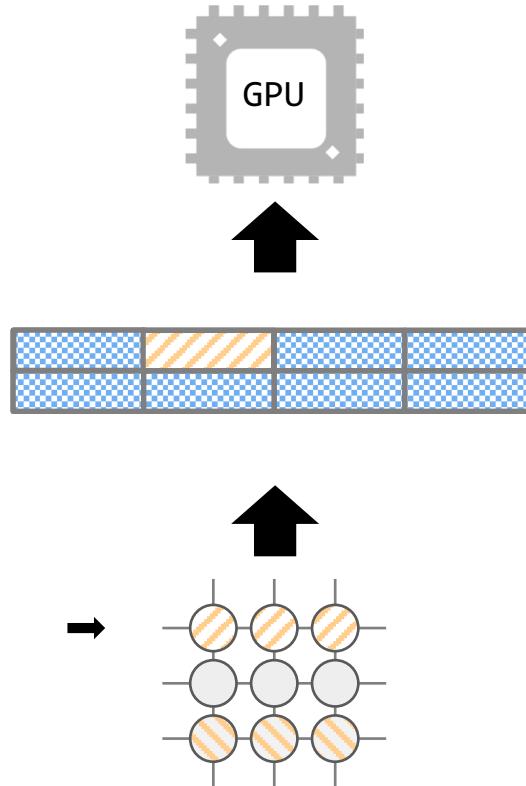


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU

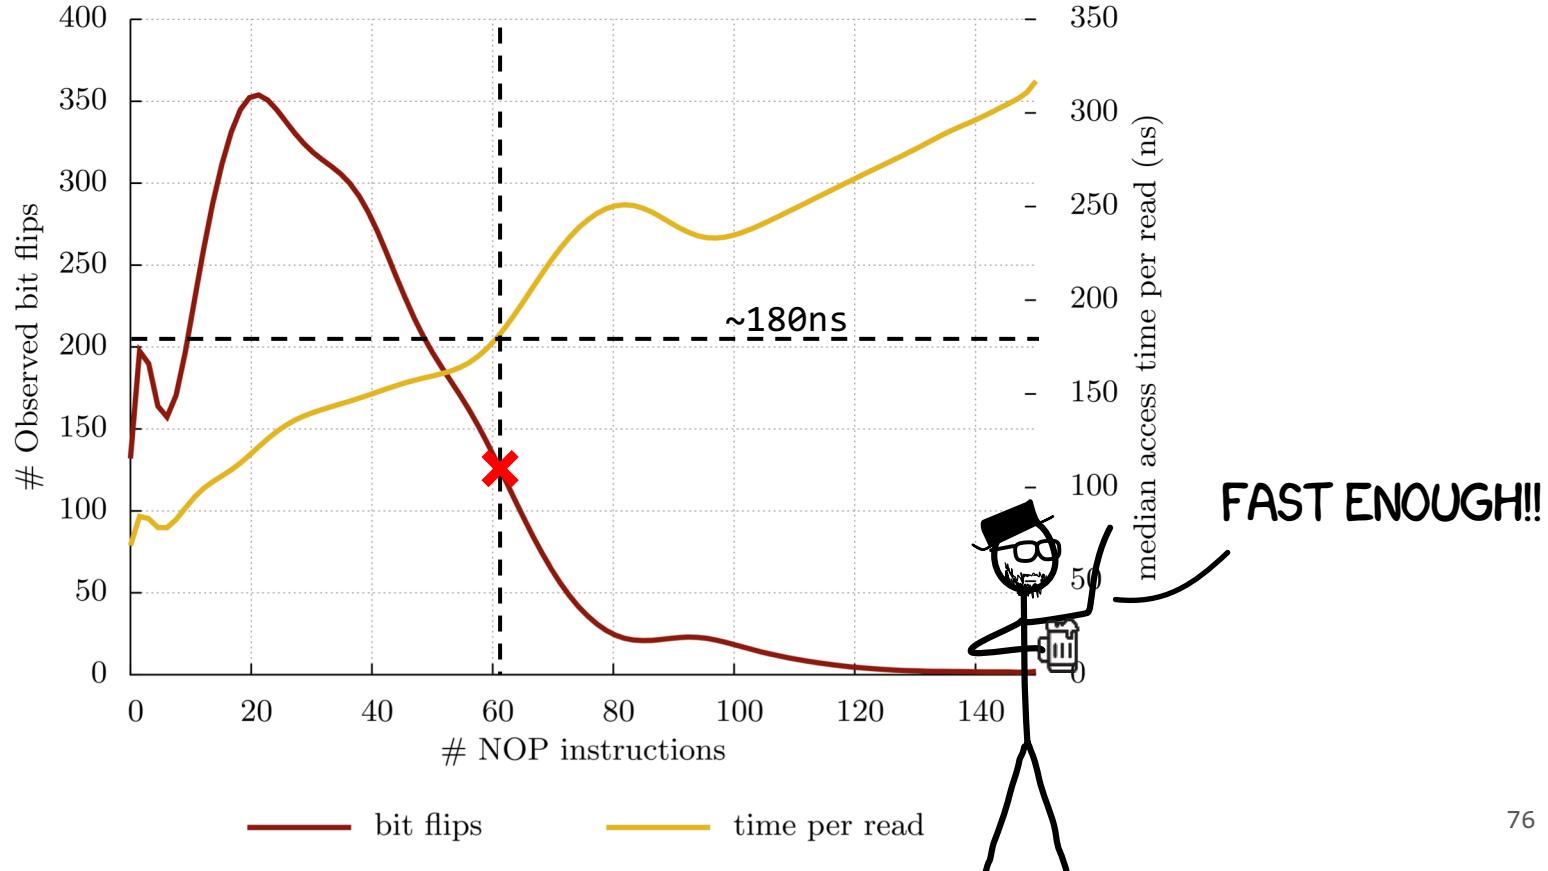


Caches: Small & Deterministic

Steps:

1. Read row $n-1$
 2. Read row $n+1$
 3. Evict++
 4. Read row $n-1$
- ...

#P2. Eviction-based Rowhammer: GPU



Attacker primitives

#P1. DRAM access ✓

#P2. Fast cache eviction ✓

#P3. Contiguous memory

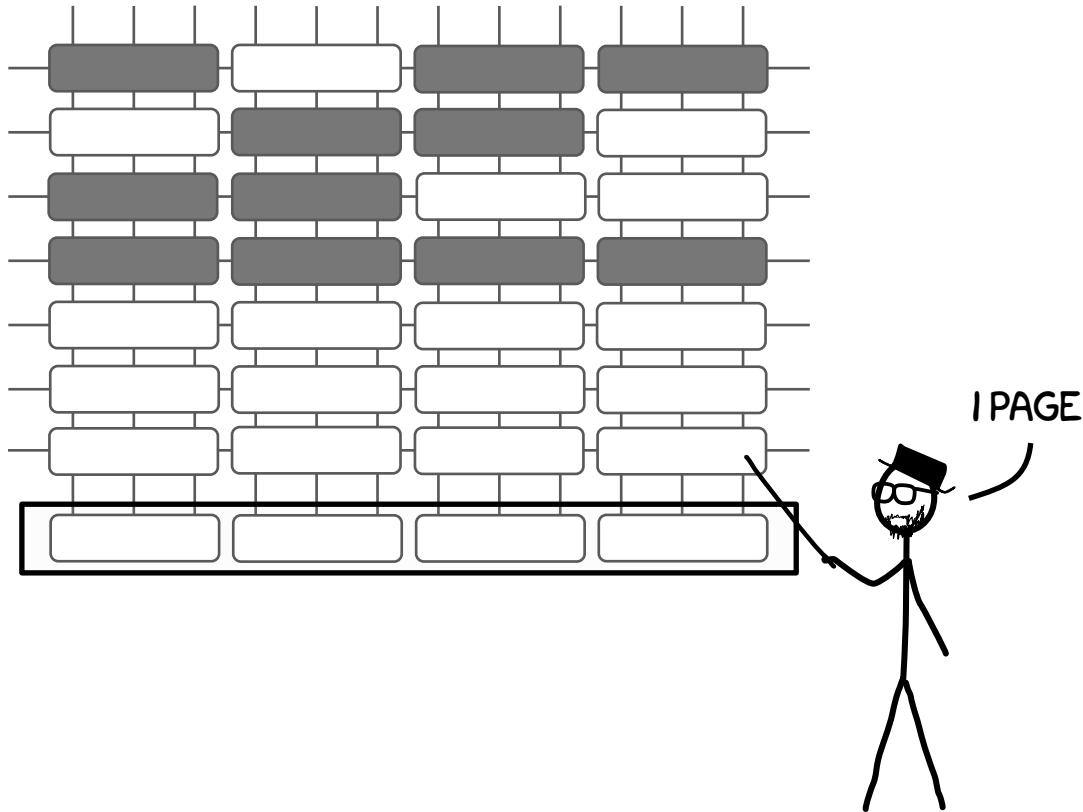
Attacker primitives

#P1. DRAM access ✓

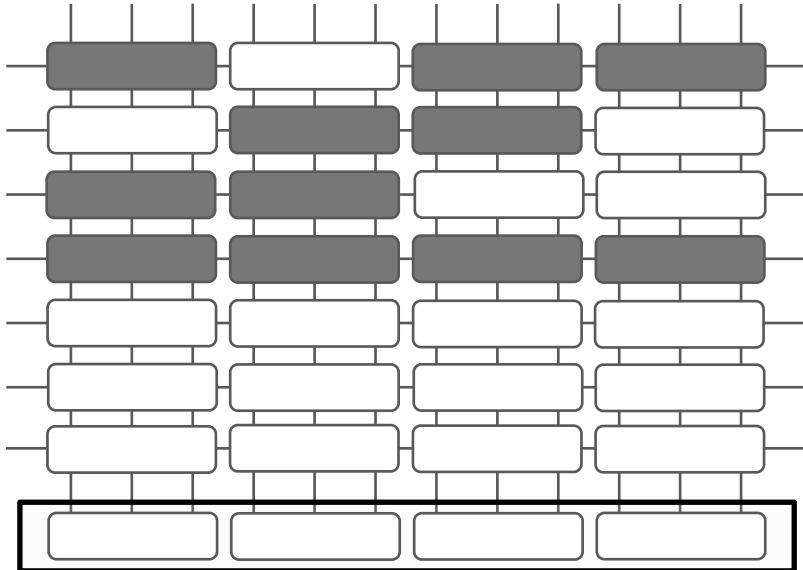
#P2. Fast cache eviction ✓

#P3. Contiguous memory

#P3. Memory Allocation

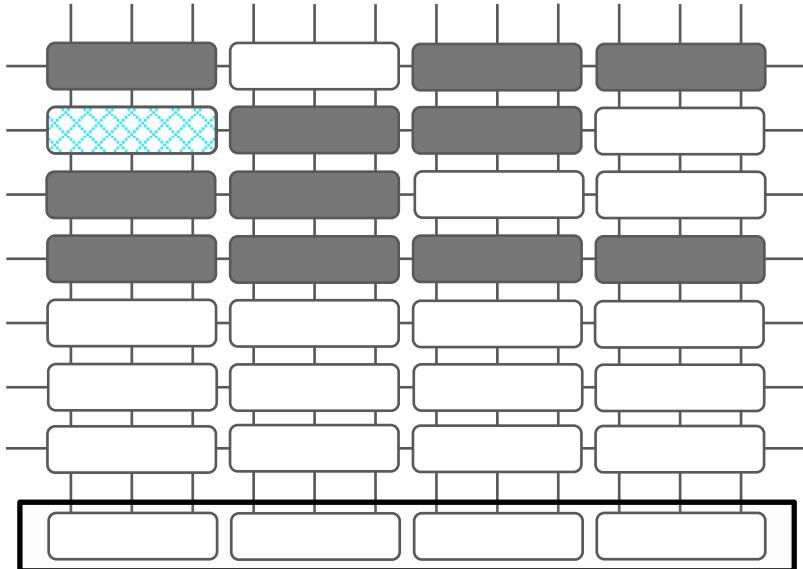


#P3. Memory Allocation



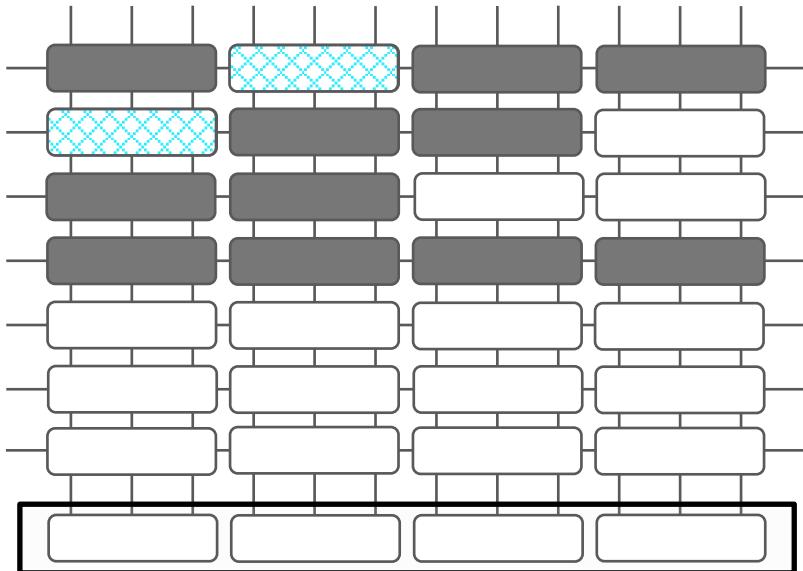
```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

#P3. Memory Allocation



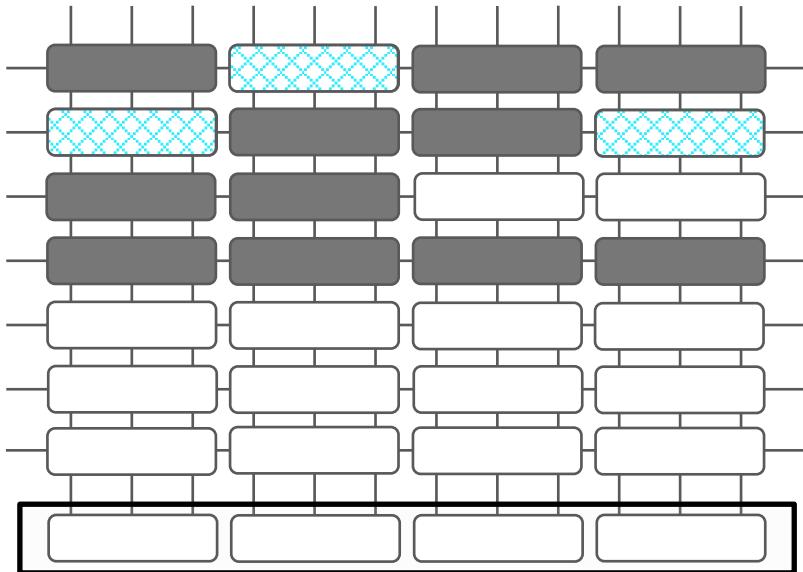
```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

#P3. Memory Allocation



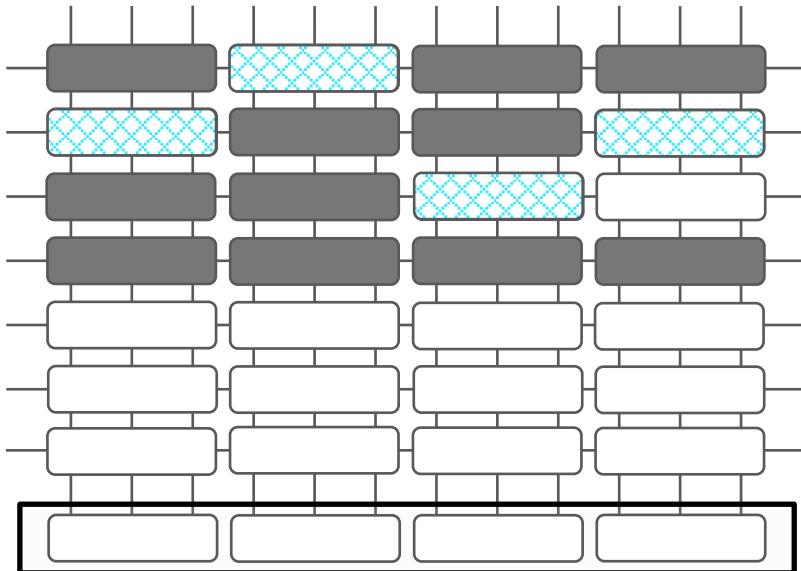
```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

#P3. Memory Allocation



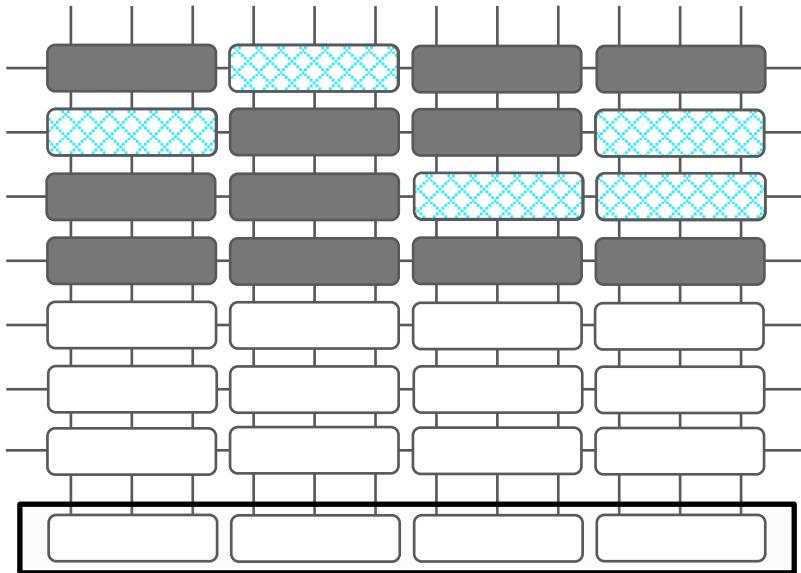
```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

#P3. Memory Allocation



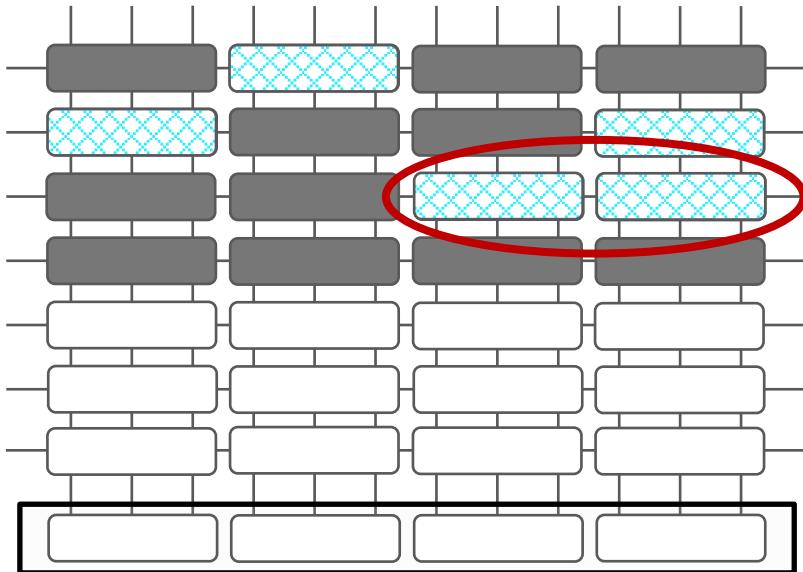
```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

#P3. Memory Allocation



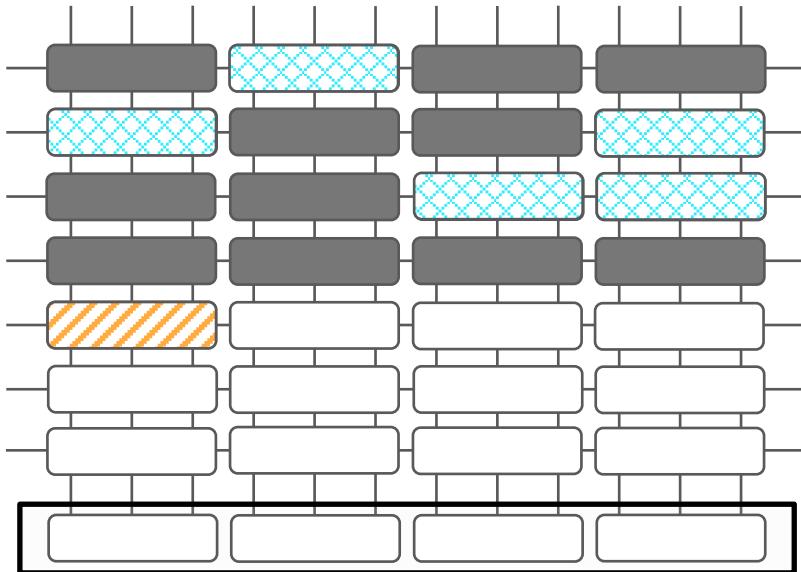
```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

#P3. Memory Allocation



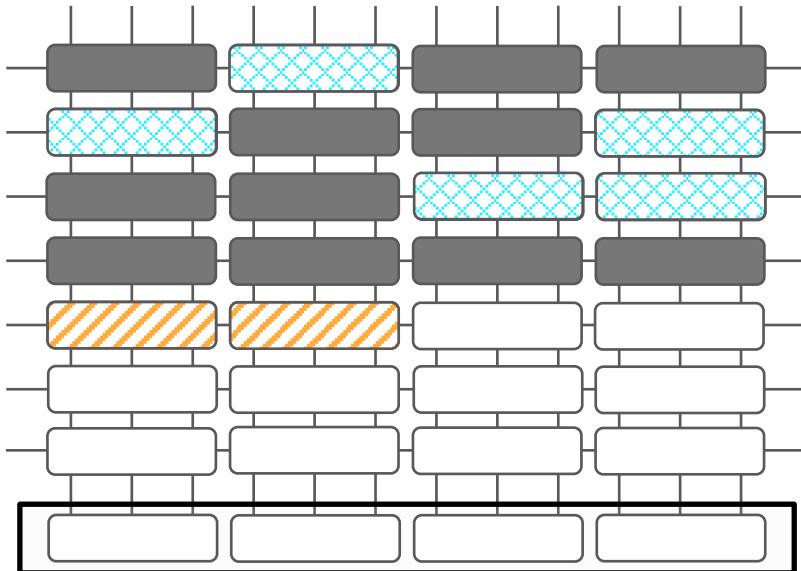
```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

#P3. Memory Allocation



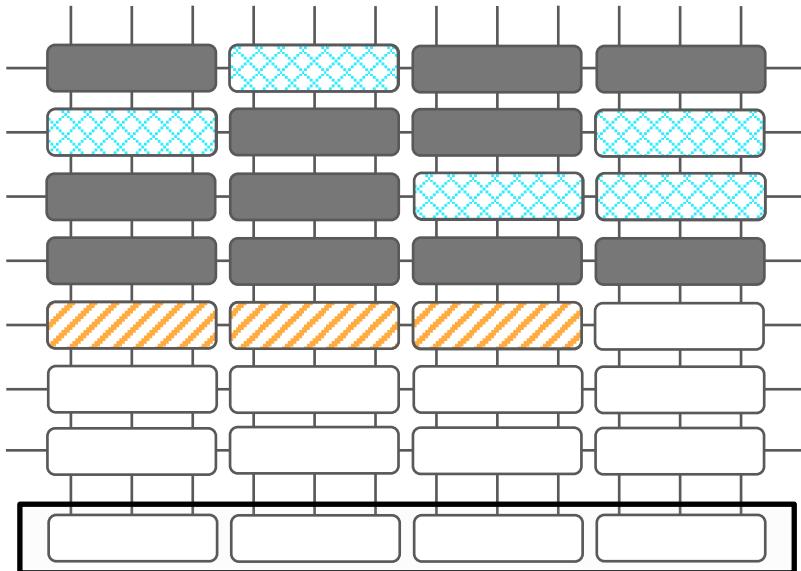
```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

#P3. Memory Allocation



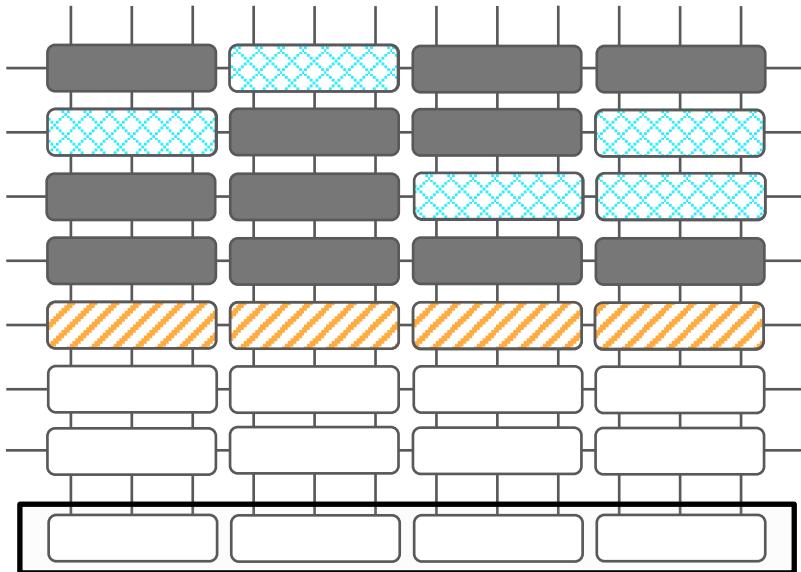
```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

#P3. Memory Allocation



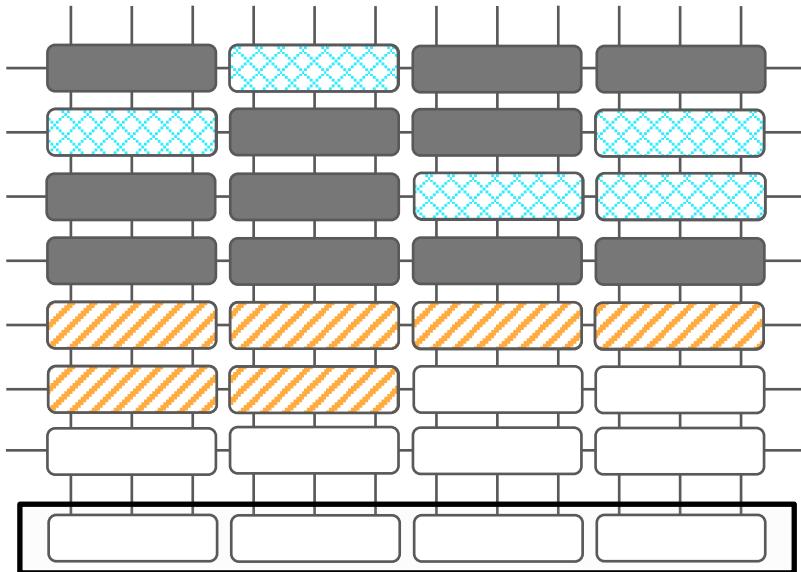
```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

#P3. Memory Allocation



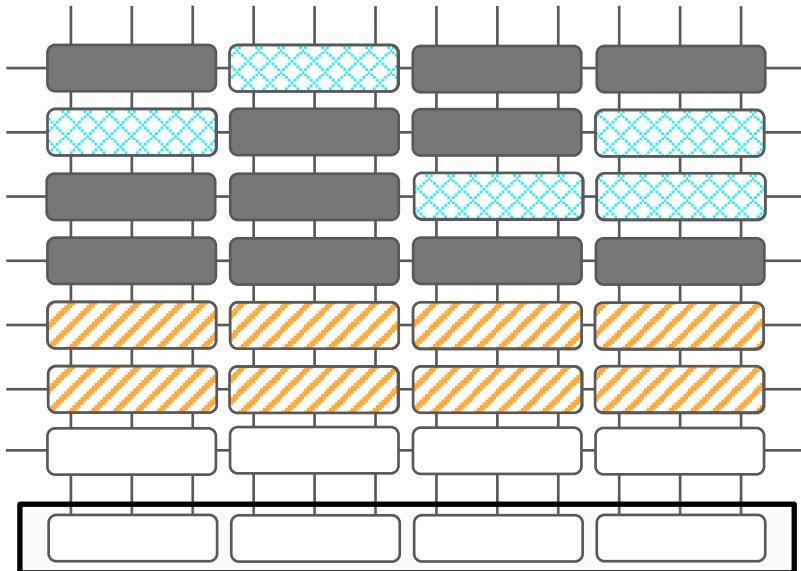
```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

#P3. Memory Allocation



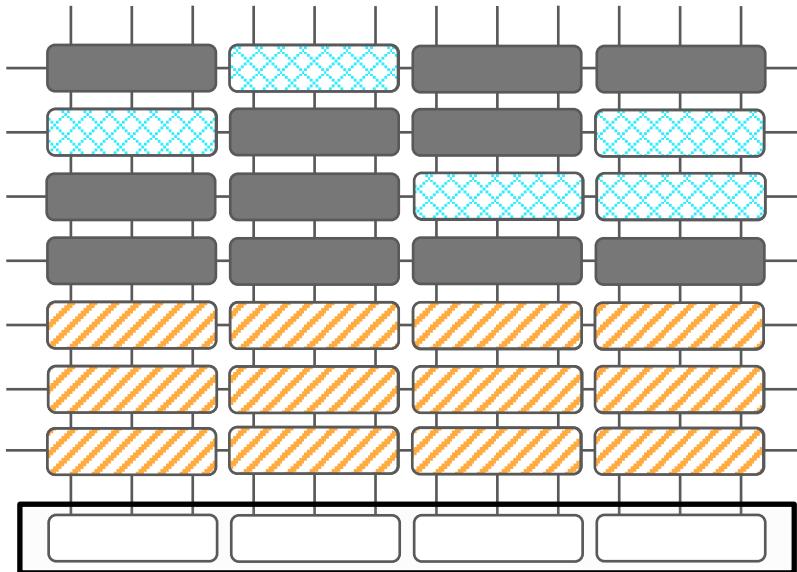
```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

#P3. Memory Allocation



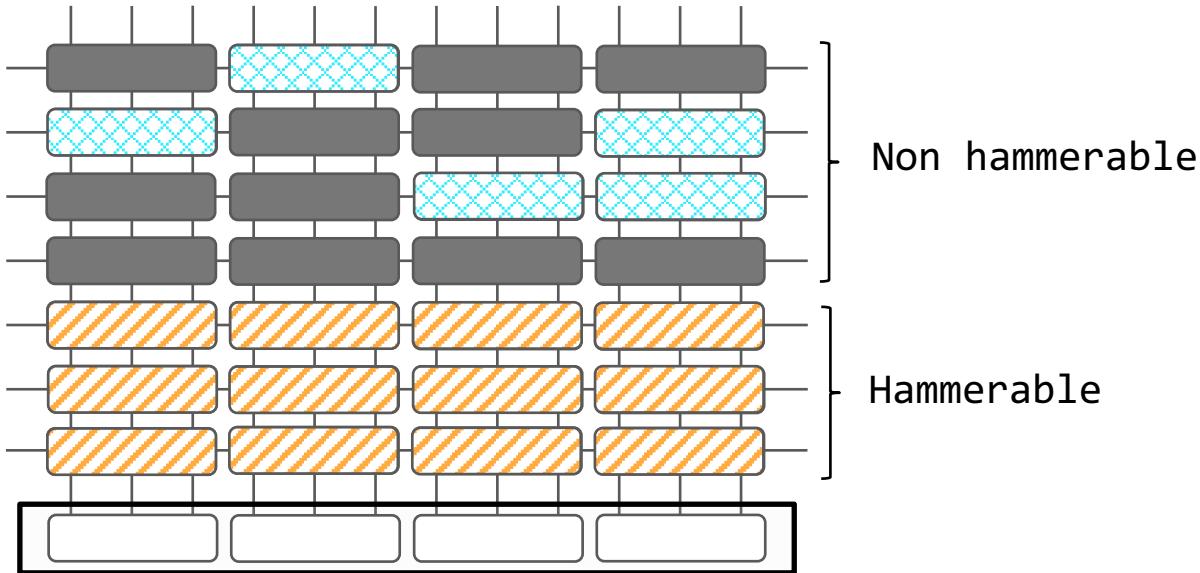
```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

#P3. Memory Allocation

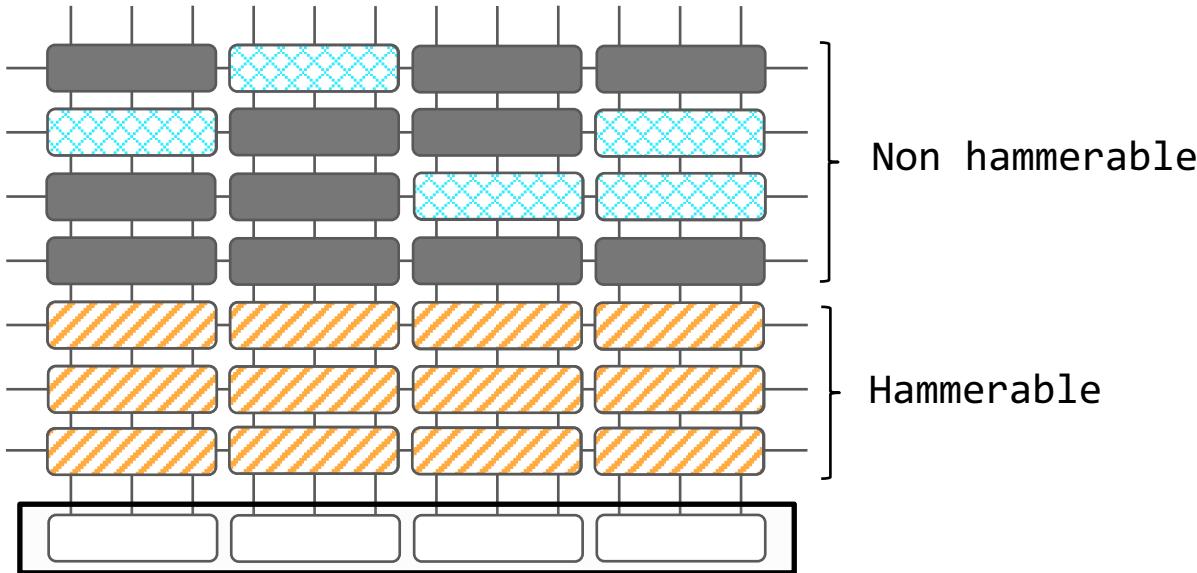


```
while (num_tex--) {  
    // size 4KB  
    tex[num_tex]= gl.createTexture();  
    fill_tex(tex[num_tex])  
}
```

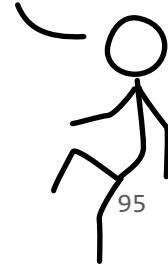
#P3. Memory Allocation



#P3. Memory Allocation



HOW DO WE
DISCERN THEM ?

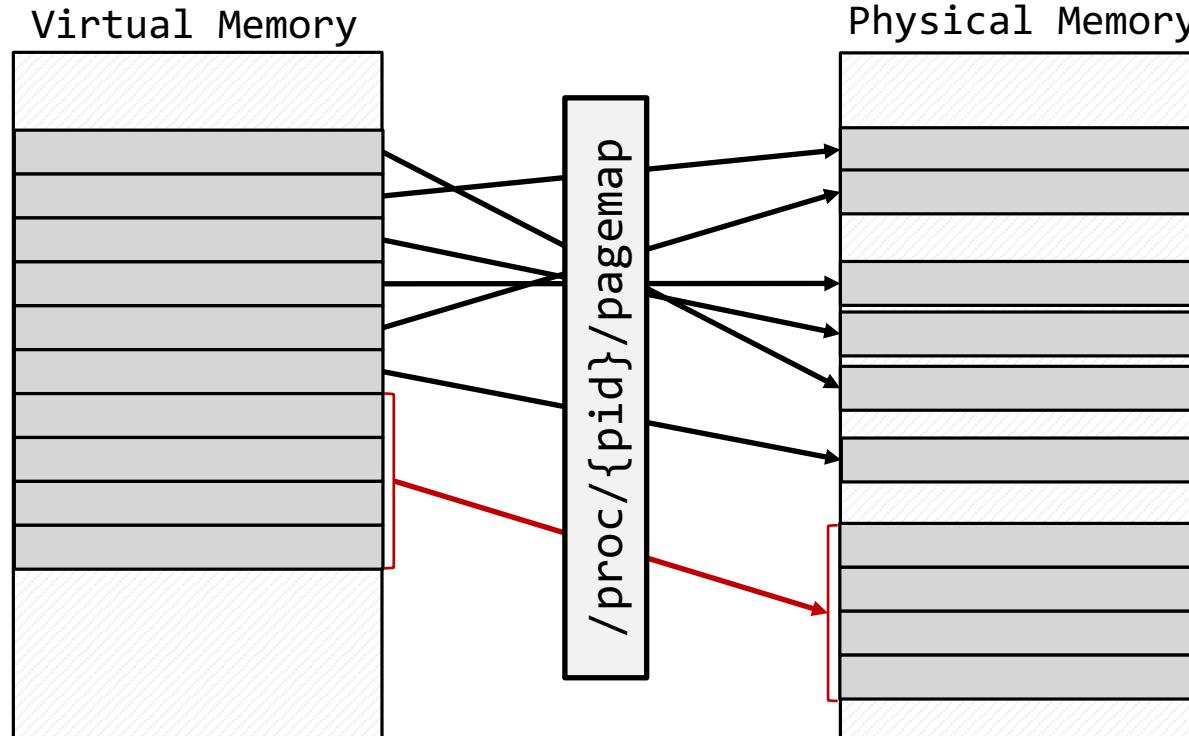


SIDE CHANNELS



SIDE CHANNELS EVERYWHERE

#P3. Contig Memory detection



Attacker primitives

#P1. DRAM access ✓

#P2. Fast cache eviction ✓

#P3. Contiguous memory ✓

What do we do with these primitives?



GLitch



GLitch: in a nutshell

Flip feng shui in JS:

0. Allocate/detect contiguous memory (side channel)
1. Memory templating
2. Memory massaging
3. Exploitation



GLitch: in a nutshell

Flip feng shui in JS:

0. Allocate/detect contiguous memory (side channel)
1. Memory templating
2. Memory massaging
3. Exploitation





Exploitation: JS Arrays

1
1.878e+65
*obj
0.3
*str
!

*func
False

arr[]

int
double
object
double
string

function
boolean

```
1 var arr = new Array(100);
2 arr[0] = 1 // int
3 arr[1] = 1.878e+65 // double
4 arr[3] = new Array(0x12) // object
...
...
```



Exploitation: JS Arrays

1
1.878e+65
*obj
0.3
*str

!

*func
False

arr[]

int
double
object
double
string

function
boolean

```
1 var arr = new Array(100);
2 arr[0] = 1 // int
3 arr[1] = 1.878e+65 // double
4 arr[3] = new Array(0x12) // object
...
...
```

NaN-boxing

IEEE-754 floating point (double)

$$1.125 = 1125 * 10^{-3}$$

$$11.25 = 1125 * 10^{-2}$$

↑ ↑
Significand exp

IEEE-754 floating point (double)

$$1.125 = 1125 * 10^{-3}$$

== !=

$$11.25 = 1125 * 10^{-2}$$

↑ ↑
Significand exp

IEEE-754 floating point (double)

$$1.125 = 1125 * 10^{-3}$$

! ==

$$11.25 = 1125 * 10^{-2}$$

IEEE-754 floating point (double)

$$1.125 = 1125 * 10^{-3}$$

== ! ==

$$11.25 = 1125 * 10^{-2}$$

↑
sign

IEEE-754 floating point (double)

$$1.125 = 1125 * 10^{-3}$$

! ==

$$11.25 = 1125 * 10^{-2}$$

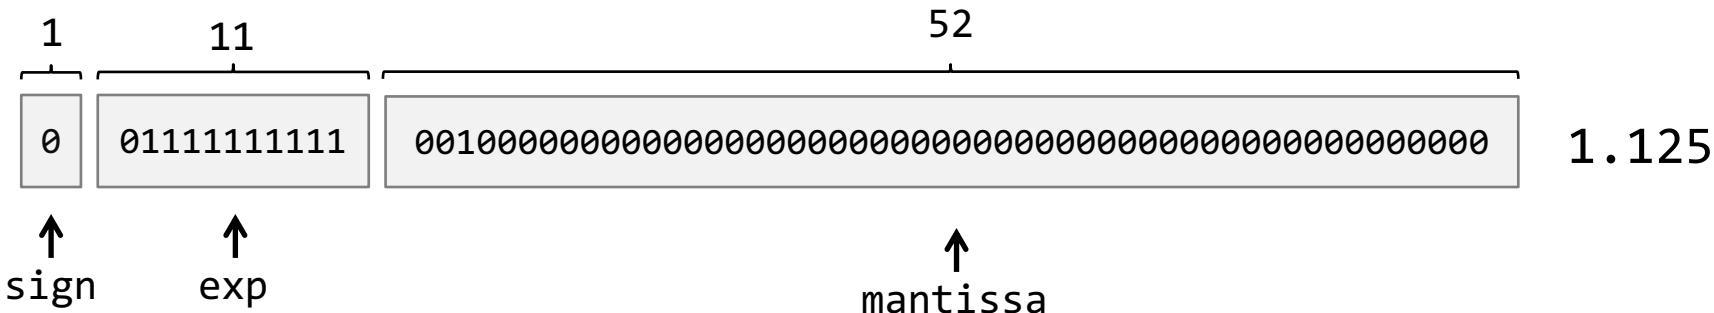


IEEE-754 floating point (double)

$$1.125 = 1125 * 10^{-3}$$

! ==

$$11.25 = 1125 * 10^{-2}$$

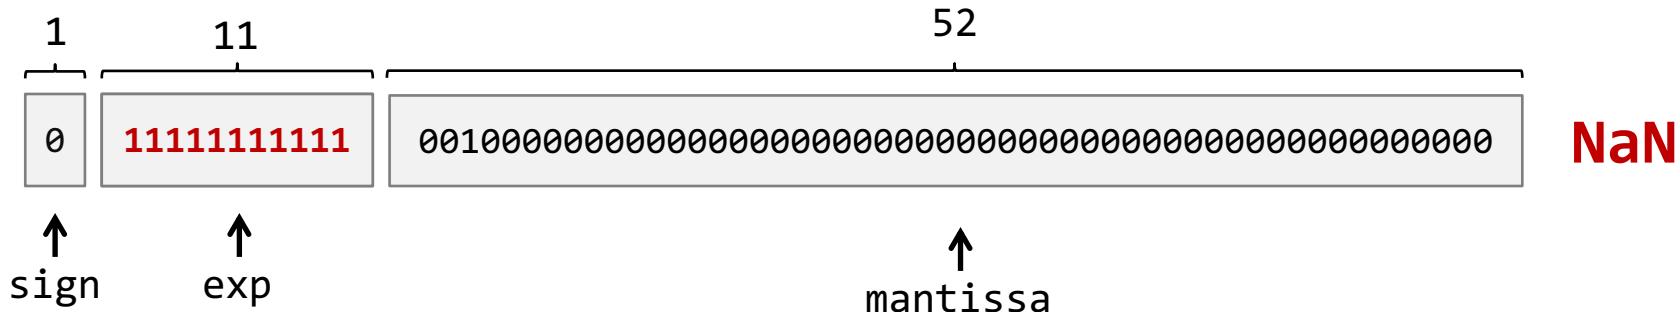


IEEE-754 floating point (double)

$$1.125 = 1125 * 10^{-3}$$

== ! ==

$$11.25 = 1125 * 10^{-2}$$

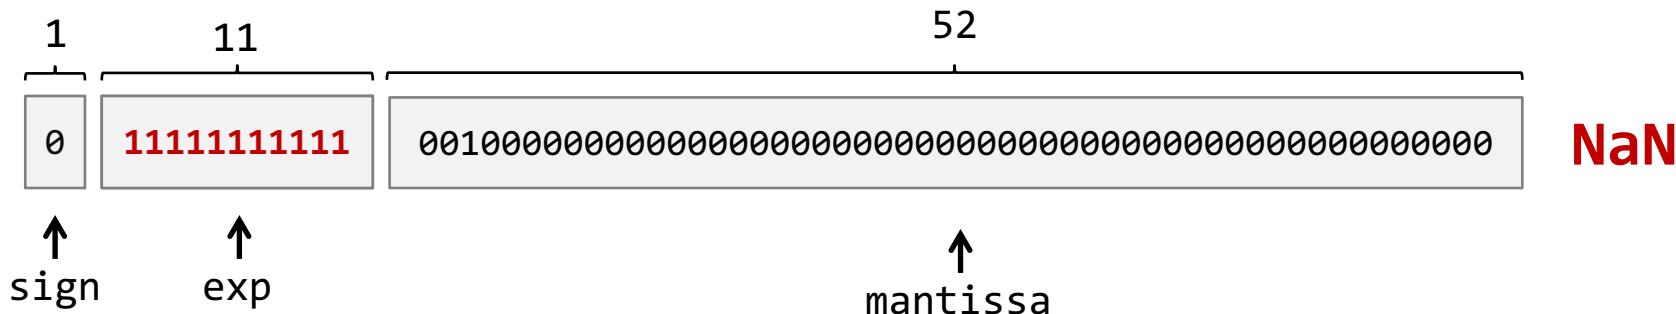


IEEE-754 floating point (double)

$$1.125 = 1125 * 10^{-3}$$

! =

$$11.25 = 1125 * 10^{-2}$$



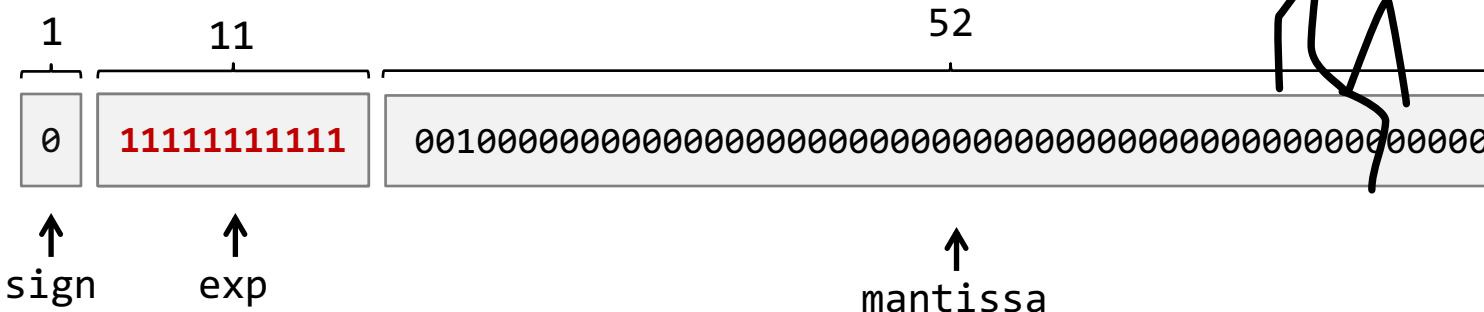
2⁵³-1 unused values

IEEE-754 floating point (double)

$$1.125 = 1125 * 10^{-3}$$

— 1 —

$$11.25 = 1125 * 10^{-2}$$



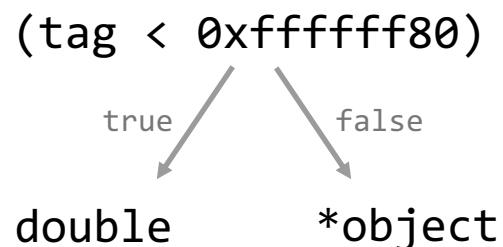
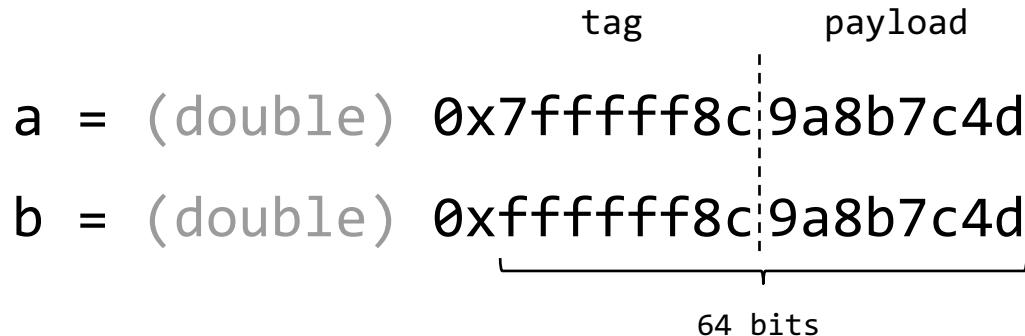
WHAT IF WE DROPOUNTERS?

NaN

2⁵³-1 unused values



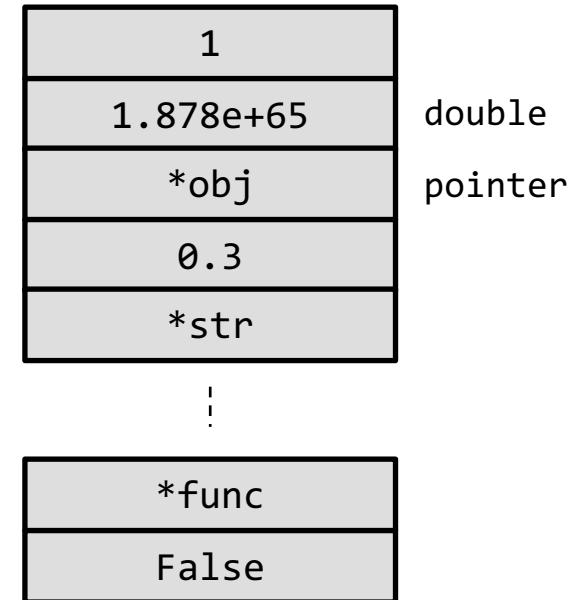
Exploitation: NuN-boxing (32bit)





Exploitation: NuN-boxing (32bit)

```
arr[1] = 0x7fffff8c|9a8b7c4d  
arr[2] = 0xfffffff8c|9a8b7c4d
```



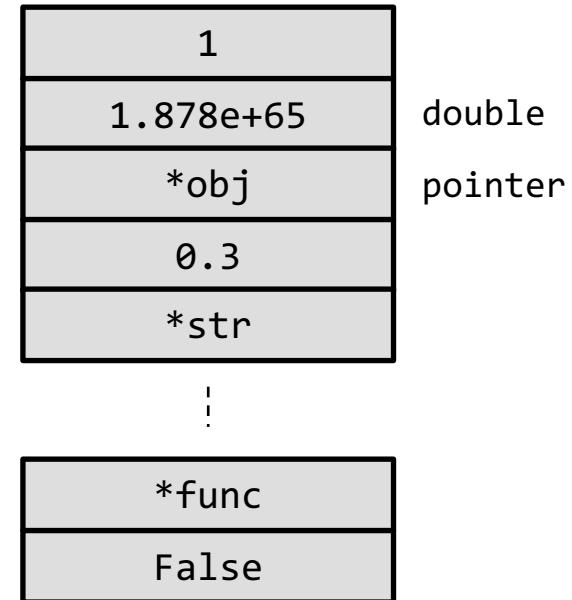


Exploitation: NuN-boxing (32bit)

arr[1] = 0x7fffff8c|9a8b7c4d
arr[2] = 0x~~ffff~~ff8c|9a8b7c4d



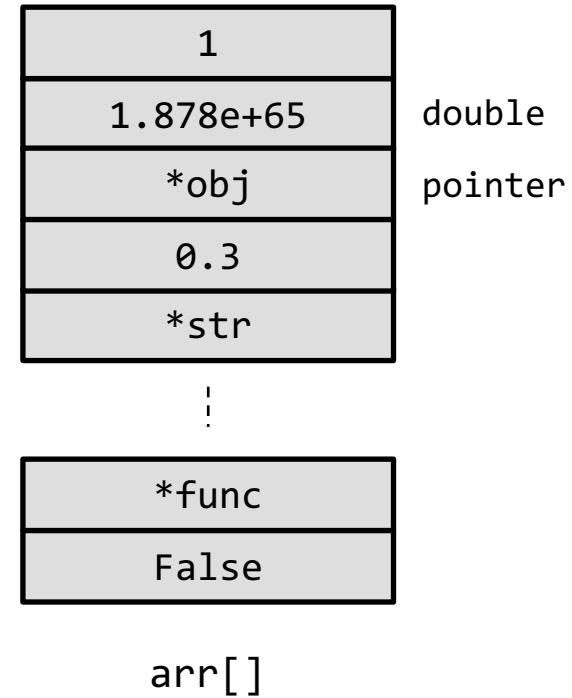
SAME PAYLOAD
1-BIT DIFFERENCE IN TAG





Exploitation: Type Flipping

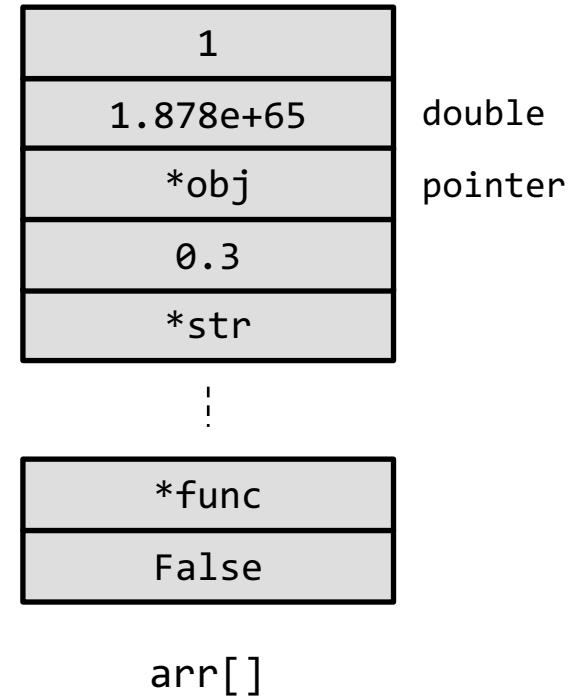
```
arr[1] = 0x7fffff8c|9a8b7c4d  
arr[2] = 0xfffffff8c|9a8b7c4d
```





Exploitation: Type Flipping

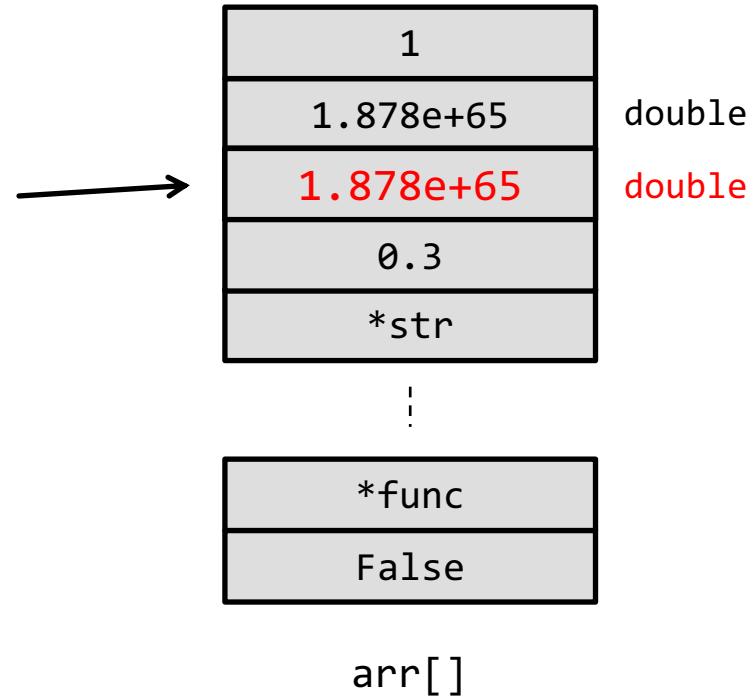
```
arr[1] = 0x7fffff8c|9a8b7c4d  
arr[2] = 0x7fffff8c|9a8b7c4d
```





Exploitation: Type Flipping

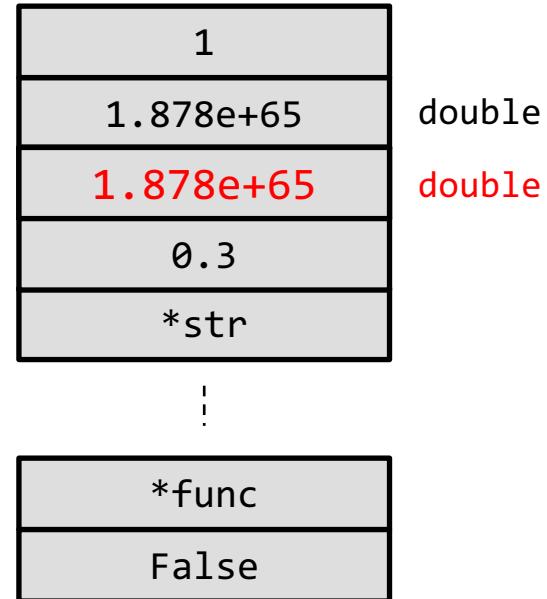
```
arr[1] = 0x7fffff8c|9a8b7c4d  
arr[2] = 0x7fffff8c|9a8b7c4d
```





Exploitation: Type Flipping

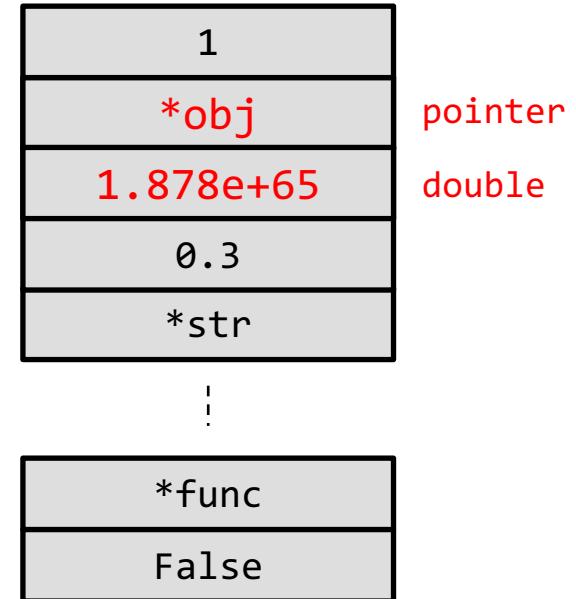
```
arr[1] = 0xfffffff8c|9a8b7c4d  
arr[2] = 0x7fffff8c|9a8b7c4d
```





Exploitation: Type Flipping

```
arr[1] = 0xfffffff8c|9a8b7c4d  
arr[2] = 0x7fffff8c|9a8b7c4d
```



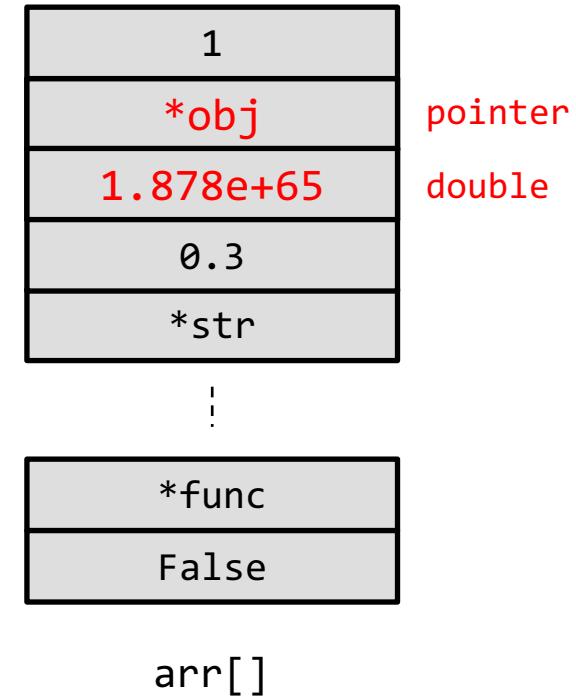


Exploitation: Type Flipping

```
arr[1] = 0xfffffff8c9a8b7c4d  
arr[2] = 0x7fffff8c9a8b7c4d
```

2 Primitives:

- #1 Arbitrary Leak [1-to-0]
- #2 Arbitrary Craft [0-to-1]



GLitch: in a nutshell



Flip feng shui in JS:

0. Allocate/detect contiguous memory (side channel)
1. Memory templating
2. Memory massaging
3. Exploitation (Arbitrary R/W)
 1. Break ASLR (1-to-0 flip)
 2. Reference fake object (0-to-1)

RUNS IN ~116 s
ON AVERAGE





WHAT ABOUT YOU?

Assignment

Steps:

- Learn about OpenGL
- Learn about GPU performance counters
- Trigger bit flips natively
- Port bit flips to JavaScript
- Exploit (Break ASLR)

} One week